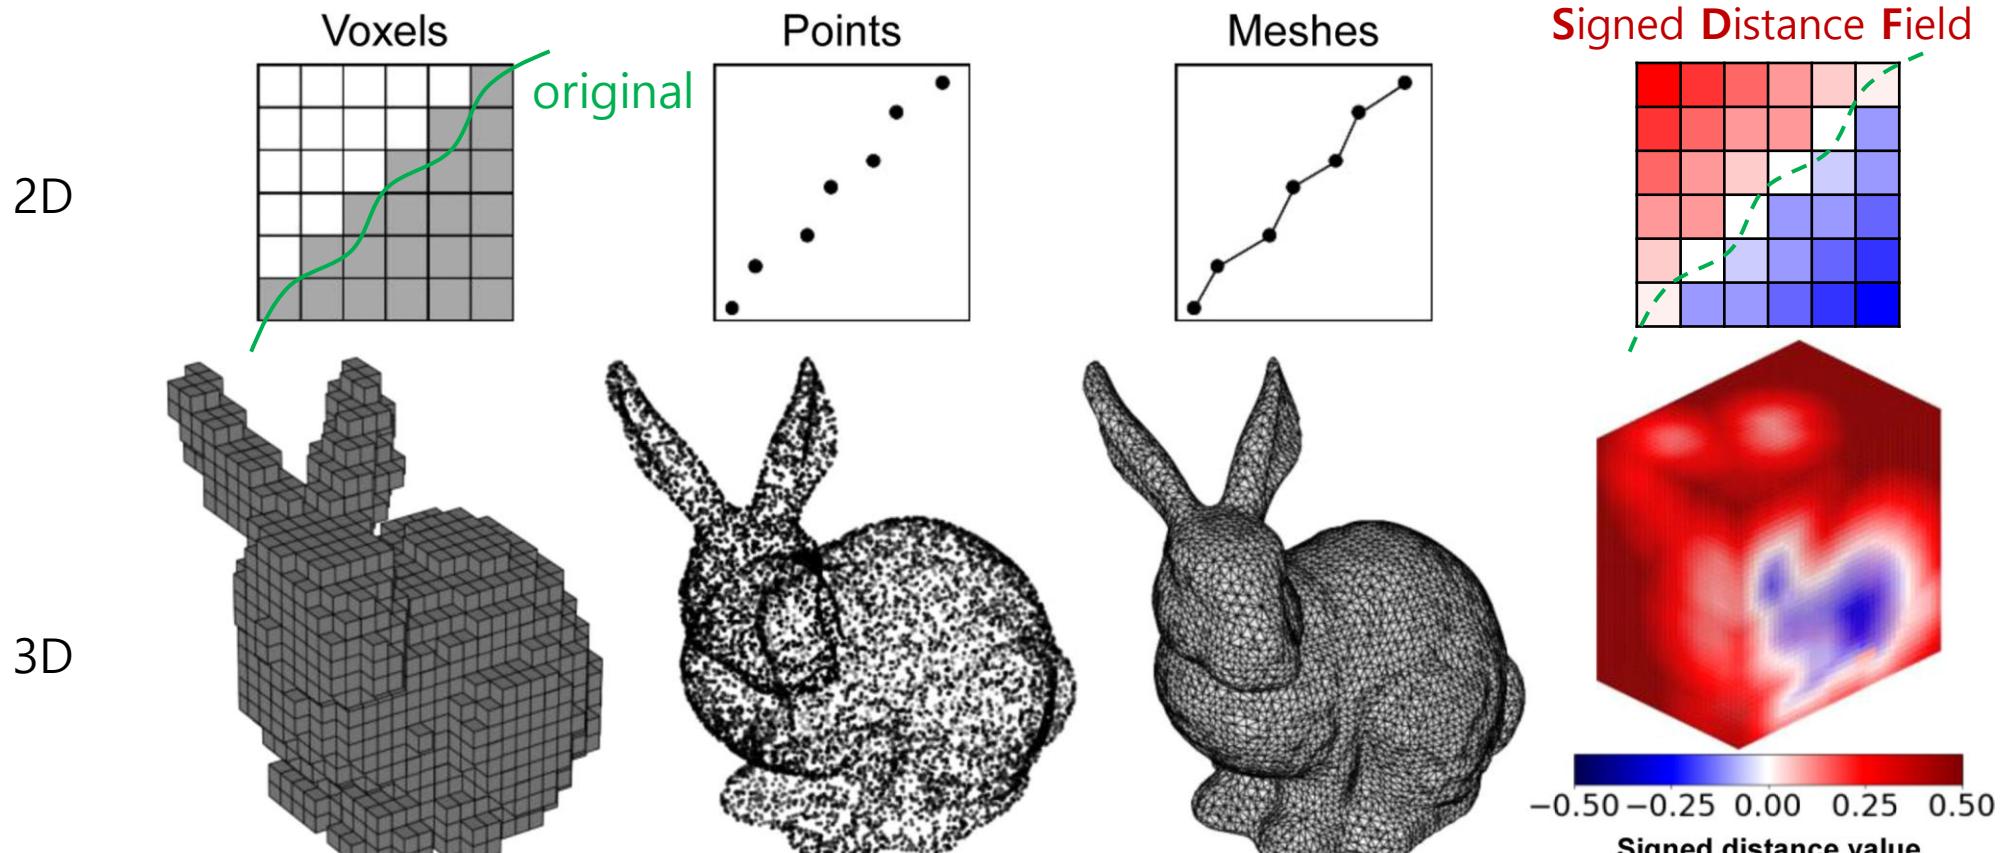


# An Invitation to 3D Vision: 3D Representations

Sunglok Choi, Assistant Professor, Ph.D.  
Dept. of Computer Science and Engineering, SEOULTECH  
[sunglok@seoultech.ac.kr](mailto:sunglok@seoultech.ac.kr) | <https://mint-lab.github.io/>

# Motivation) 3D Shape Representations

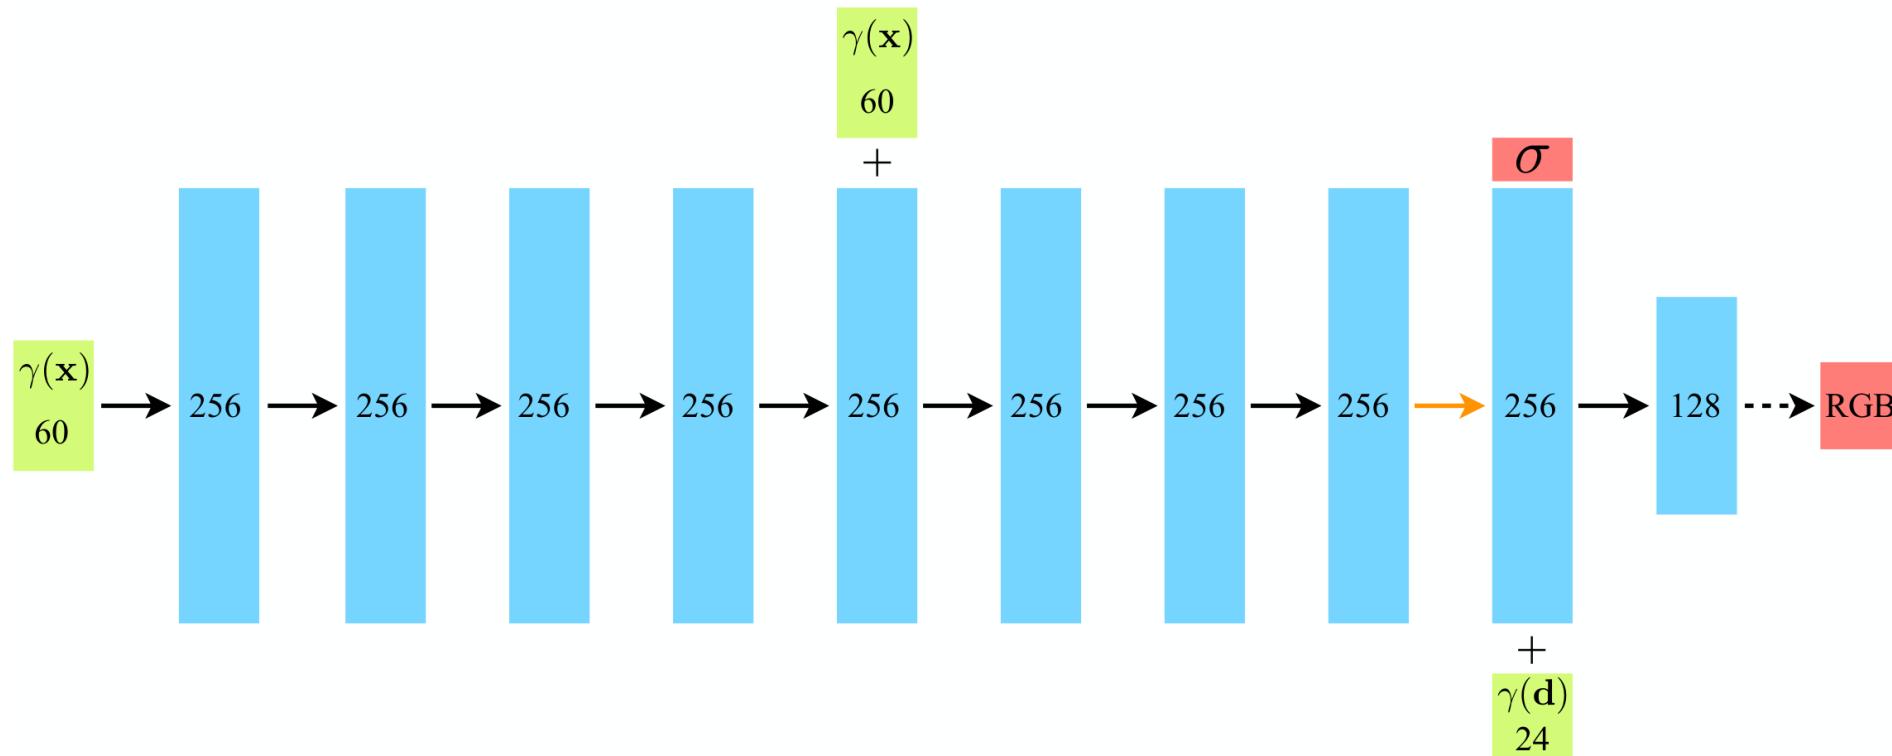
- Explicit vs. **Implicit** representations



Why not a **neural network**?

# NeRF (Neural Radiance Field; 2020)

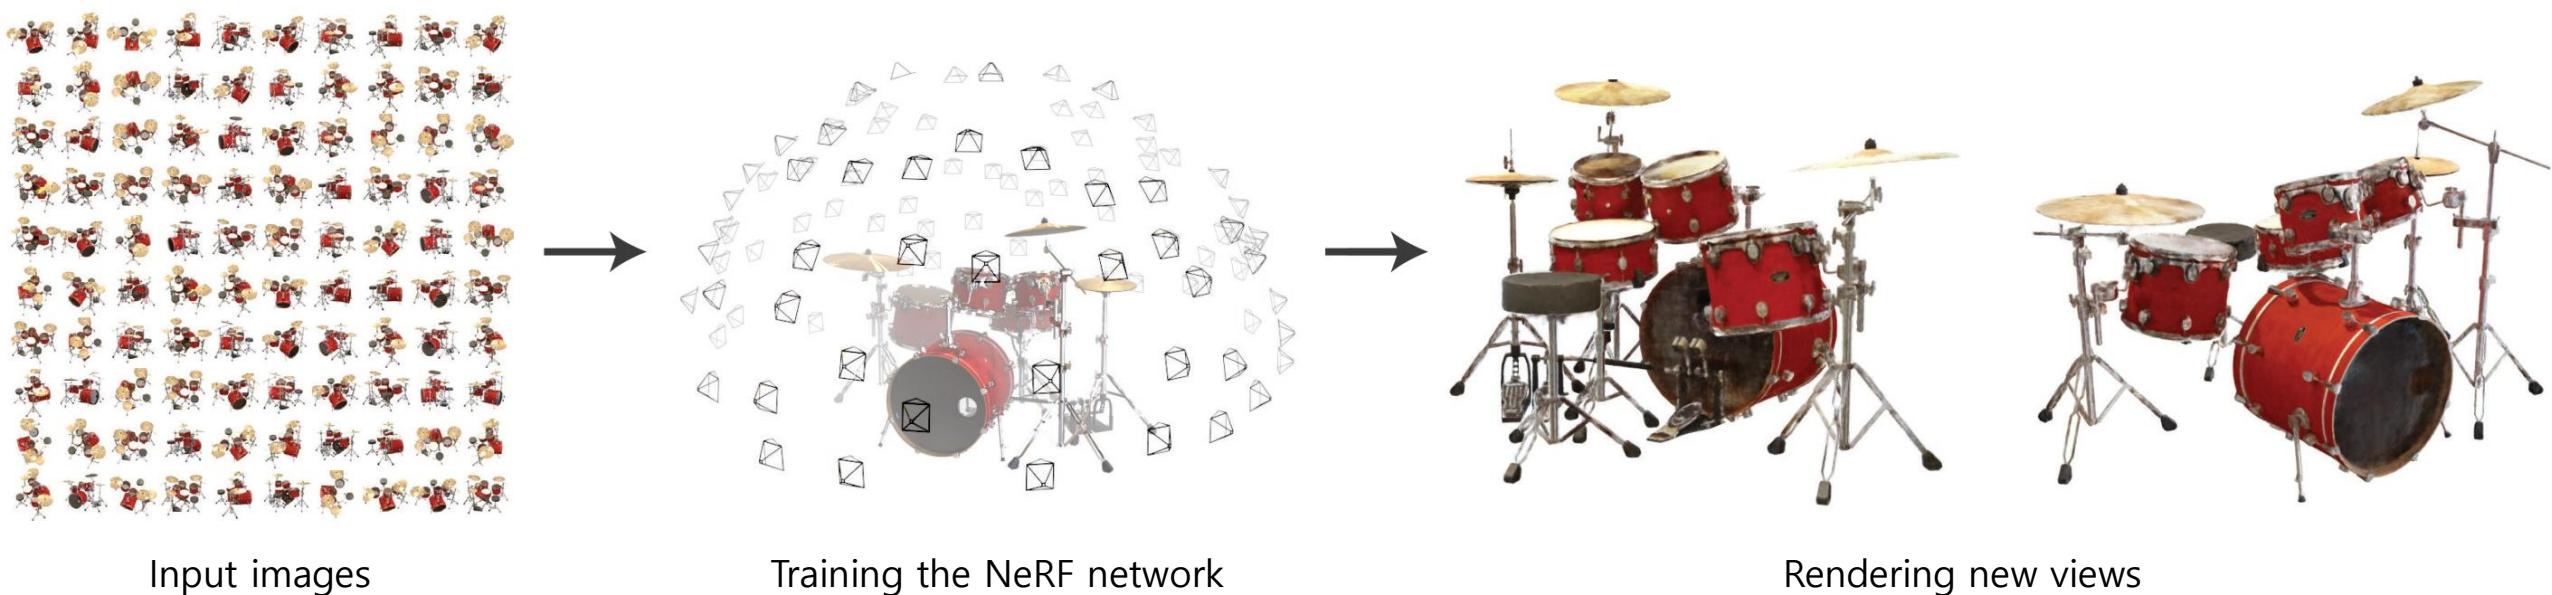
- NeRF is a **multi-layer perceptron** for rendering an image from a new viewpoint.
- Network: **11 fully-connected (shortly FC) layers**
  - **Input:** Spatial location  $\mathbf{x} = (x, y, z)$  and viewing direction  $\mathbf{d} = (d_x, d_y, d_z)$  on a unit sphere
  - **Output:** RGB color (RGB) and density ( $\sigma$ )



# NeRF (Neural Radiance Field; 2020)

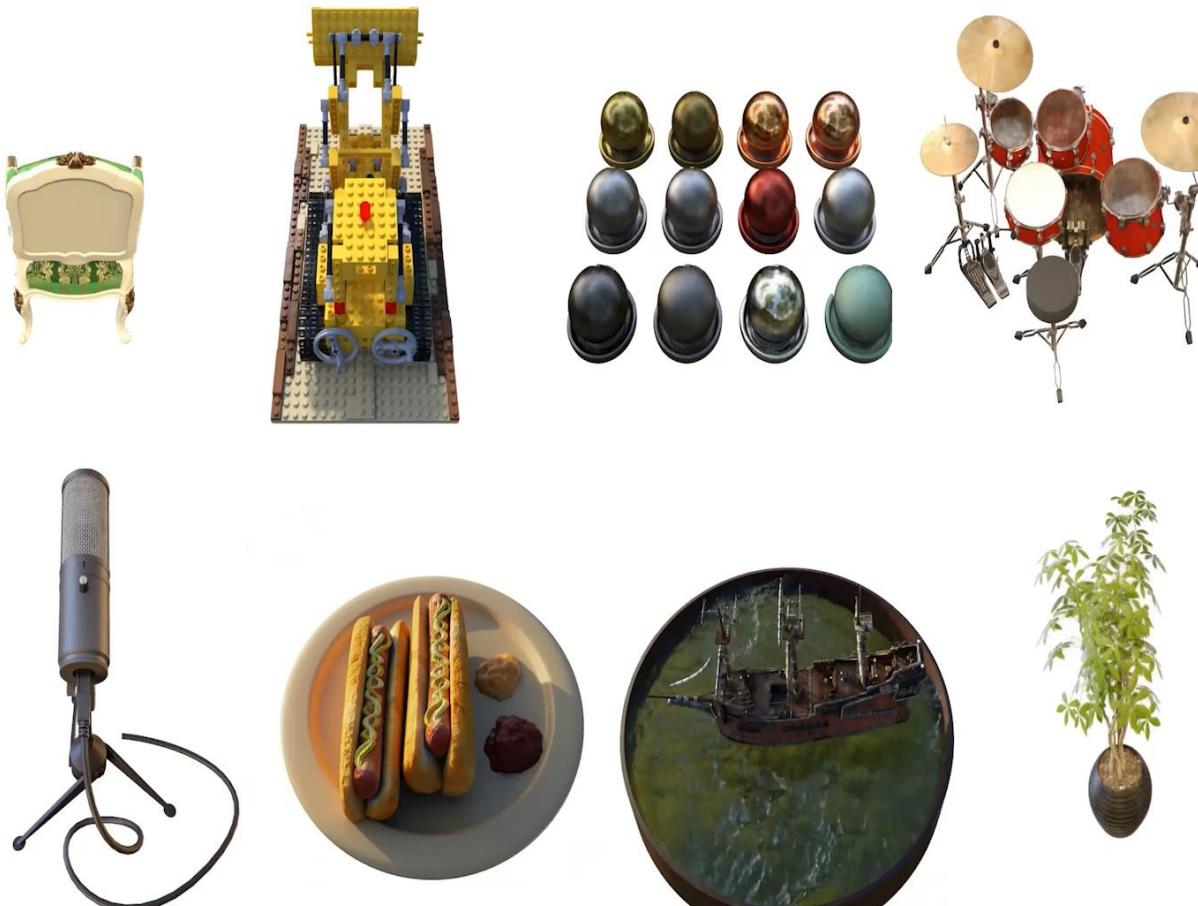
implicitly representing a 3D scene.

- NeRF is a multi-layer perceptron for ~~rendering an image from a new viewpoint~~.
- **Training:** Learning the 3D scene
  - Input: Images with their 3D viewpoints ( $R_j, t_j$ )
    - Note) 3D viewpoints can be retrieved by SfM (e.g. COLMAP).
- **Inference:** Synthesizing a 2D image with a *new* viewpoint
  - Input: A new camera viewpoint ( $R_n, t_n$ )



# NeRF (Neural Radiance Field; 2020)

- NeRF is a multi-layer perceptron for **rendering an image from a new viewpoint**.
- **Inference:** Synthesizing a 2D image with a *new* viewpoint
  - Input: A new camera viewpoint ( $R_n$ ,  $t_n$ )



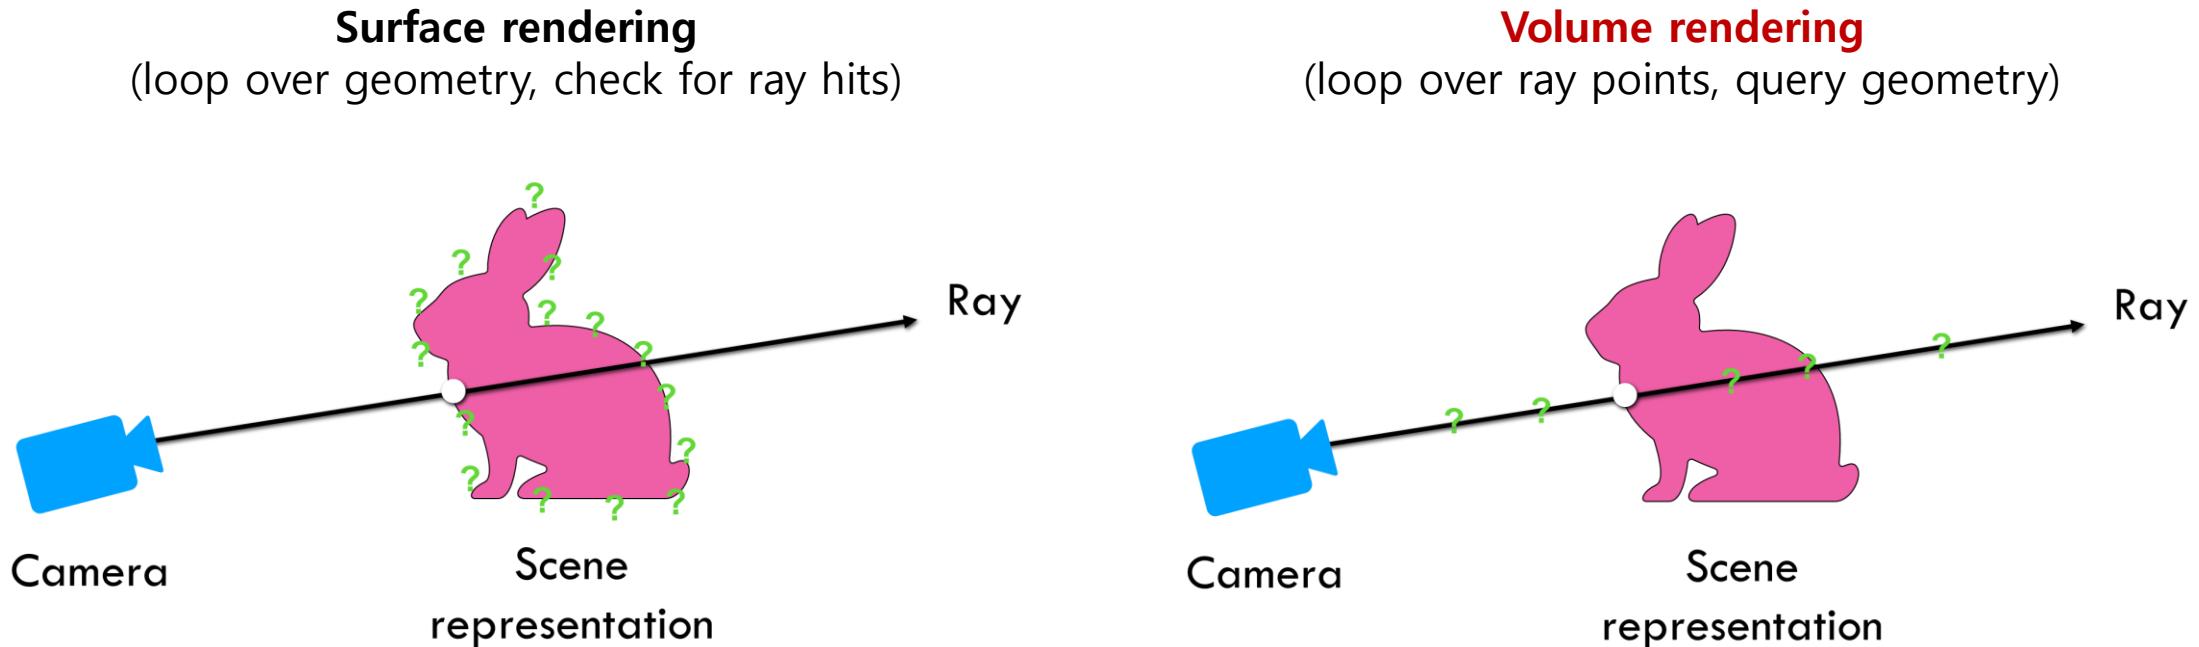
# NeRF (Neural Radiance Field; 2020)

- Key idea: **Neural Volumetric Rendering**
  - + Continuous rendering
  - + Differentiable rendering
  - + Model without concrete ray/surface intersections



# NeRF (Neural Radiance Field; 2020)

- Key idea: **Neural Volumetric Rendering**

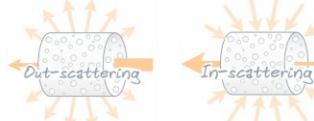


# NeRF (Neural Radiance Field; 2020)

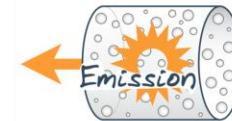
- Key idea: **Neural Volumetric Rendering**
  - Based on the simplified physics (ignoring **scattering**)



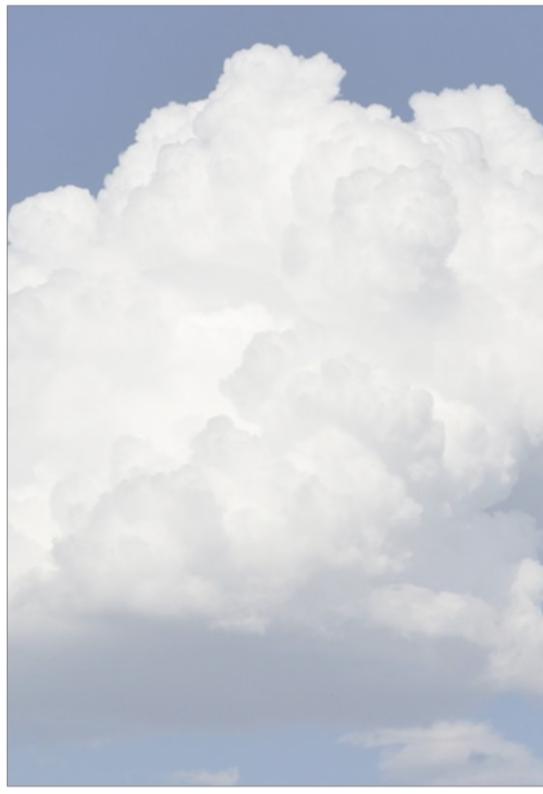
Absorption



Scattering

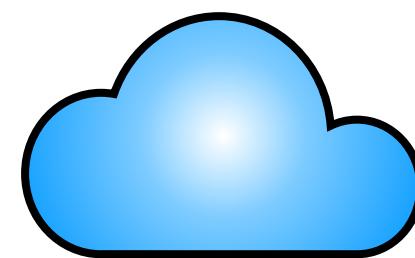


Emission



# NeRF (Neural Radiance Field; 2020)

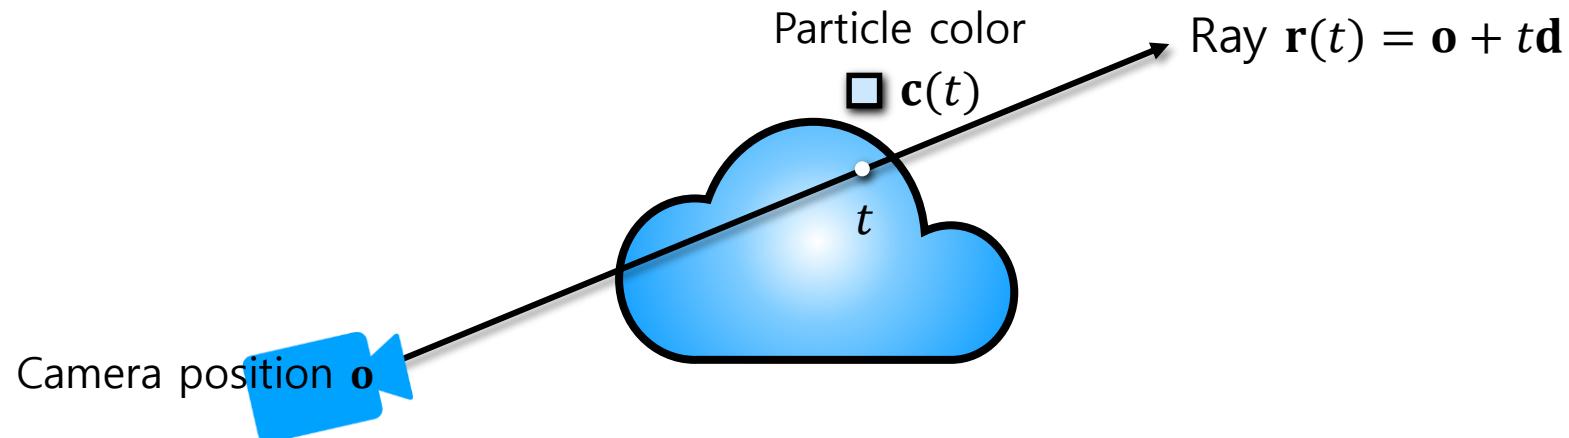
- Key idea: **Neural Volumetric Rendering**
  - The scene is a **cloud** composed of tiny colored particles.



3D volume

# NeRF (Neural Radiance Field; 2020)

- Key idea: **Neural Volumetric Rendering**
  - If a **ray** (traveling through the scene) hits a **particle** at distance  $t$  (along the ray), we can retrieve its color  $\mathbf{c}(t)$ .



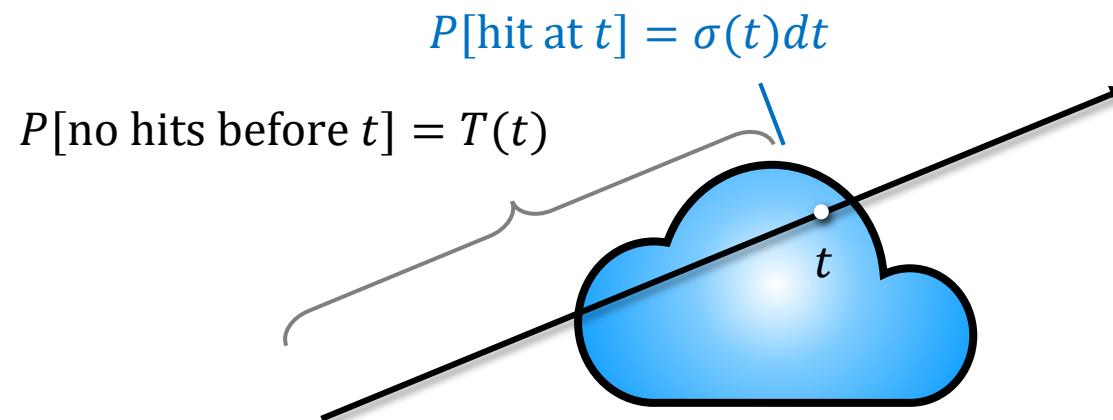
# NeRF (Neural Radiance Field; 2020)

- Key idea: **Neural Volumetric Rendering**

- The product of these probabilities tells us how much you see the particles at  $t$ :

- $P[\text{first hit at } t] = P[\text{no hit before } t] \times P[\text{hit at } t] = T(t)\sigma(t)dt$

- $T(t)$ : Transmittance (the probability that the ray doesn't hit any particles earlier)



# NeRF (Neural Radiance Field; 2020)

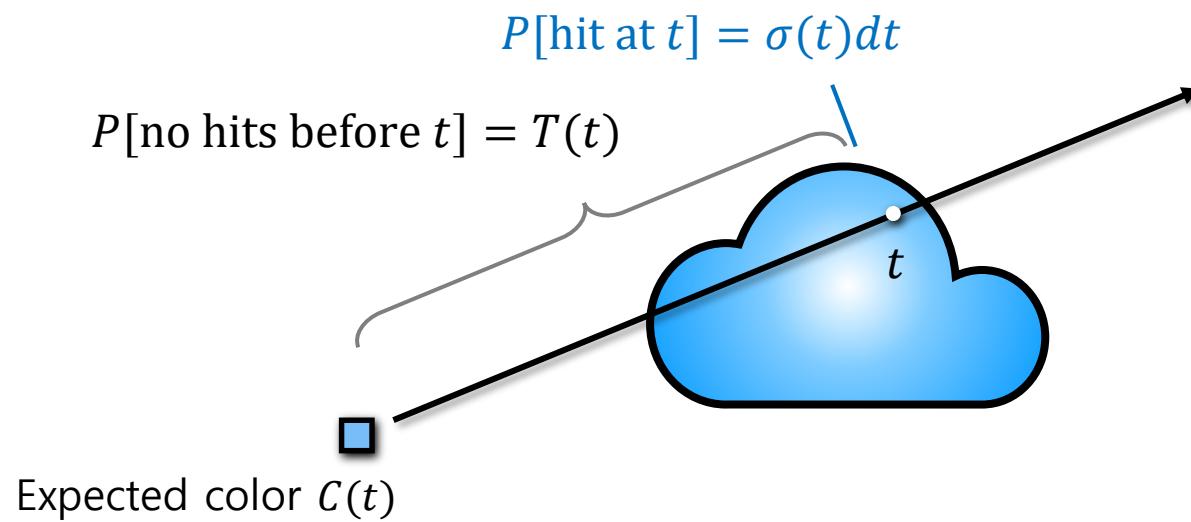
- Key idea: **Neural Volumetric Rendering**

- The product of these probabilities tells us how much you see the particles at  $t$ :

- $P[\text{first hit at } t] = P[\text{no hit before } t] \times P[\text{hit at } t] = T(t)\sigma(t)dt$

- Expected color by a ray  $\mathbf{r}$

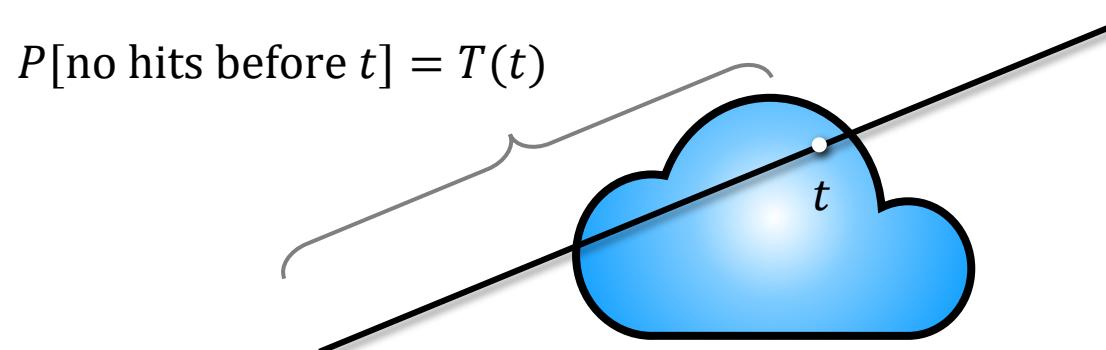
$$C(\mathbf{r}) = \int_{t_0}^{t_1} T(t)\sigma(t)\mathbf{c}(t)dt \approx \sum_{i=1}^N T_i \mathbf{c}_i (1 - \exp(-\sigma_i \delta_i)) \quad \text{where} \quad T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$



# NeRF (Neural Radiance Field; 2020)

- Key idea: **Neural Volumetric Rendering**
  - Q) Can we represent  $T(t)$  using the (known) density function  $\sigma(t)$ ?
    - A recursive form:  $P[\text{no hit before } t + dt] = P[\text{no hit before } t] \times P[\text{no hit at } t]$

$$T(t + dt) = T(t)(1 - \sigma(t)dt) \rightarrow T(t) = \exp\left(-\int_{t_0}^t \sigma(s)ds\right)$$



# NeRF (Neural Radiance Field; 2020)

- Key idea: **Neural Volumetric Rendering**

- Q) Can we represent  $T(t)$  using the (known) density function  $\sigma(t)$ ?

- A recursive form:  $P[\text{no hit before } t + dt] = P[\text{no hit before } t] \times P[\text{no hit at } t]$

$$T(t + dt) = T(t)(1 - \sigma(t)dt) \rightarrow T(t) = \exp\left(-\int_{t_0}^t \sigma(s)ds\right)$$

- Solving the differential equation

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

$$T(t) + T'(t)dt = T(t) - T(t)\sigma(t)dt \quad (\text{Taylor expansion for } T)$$

$$\frac{T'(t)}{T(t)}dt = -\sigma(t)dt \quad (\text{Rearrange})$$

$$\log T(t) = -\int_{t_0}^t \sigma(s)ds \quad (\text{Integrate from } t_0 \text{ to } t)$$

$$T(t) = \exp\left(-\int_{t_0}^t \sigma(s)ds\right) \quad (\text{Exponentiate})$$

# NeRF (Neural Radiance Field; 2020)

- Key idea: **Neural Volumetric Rendering**

- Q) How can we integrate the color equation?

$$C(\mathbf{r}) = \int_{t_0}^{t_1} T(t) \sigma(t) \mathbf{c}(t) dt \approx \sum_{i=1}^N T_i \mathbf{c}_i (1 - \exp(-\sigma_i \delta_i))$$

- Using the quadrature rule (구분구적법 in Korean)

$$\int T(t) \sigma(t) \mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t) \sigma_i \mathbf{c}_i dt \quad (\text{Quadrature rule})$$

$$= \sum_{i=1}^n T_i \sigma_i \mathbf{c}_i \int_{t_i}^{t_{i+1}} \exp(-\sigma_i(t - t_i)) dt \quad (\text{Substitute})$$

$$= \sum_{i=1}^n T_i \sigma_i \mathbf{c}_i \frac{\exp(-\sigma_i(t_{i+1} - t_i)) - 1}{-\sigma_i} \quad (\text{Integrate})$$

$$= \sum_{i=1}^n T_i \mathbf{c}_i (1 - \exp(-\sigma_i \delta_i)) \quad (\text{Cancel } \sigma_i \text{ and substitute } \delta_i = t_{i+1} - t_i)$$

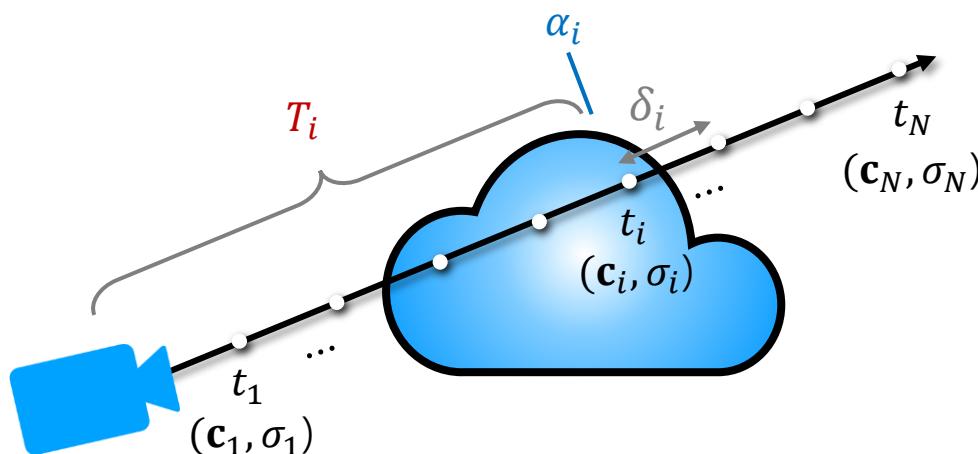
# NeRF (Neural Radiance Field; 2020)

- Key idea: **Neural Volumetric Rendering**

- Rendering an image for a ray  $\mathbf{r} = \mathbf{o} + t\mathbf{d}$

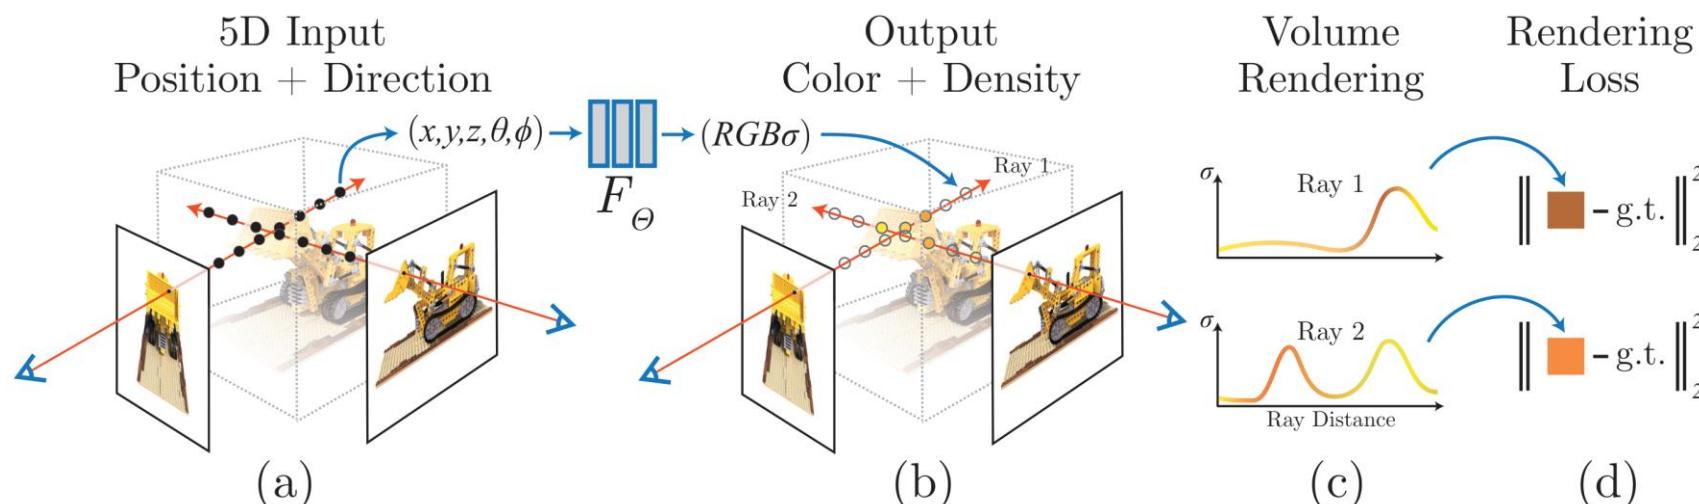
$$C(\mathbf{r}) = \sum_{i=1}^N T_i \alpha_i \mathbf{c}_i \quad \text{where} \quad \alpha_i = 1 - \exp(-\sigma_i \delta_i) \quad \text{and} \quad T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

- $T_i$ : How much light is blocked before ray segment  $i$ ?
  - $\alpha_i$ : How much light is contributed by ray segment  $i$ ?



# NeRF (Neural Radiance Field; 2020)

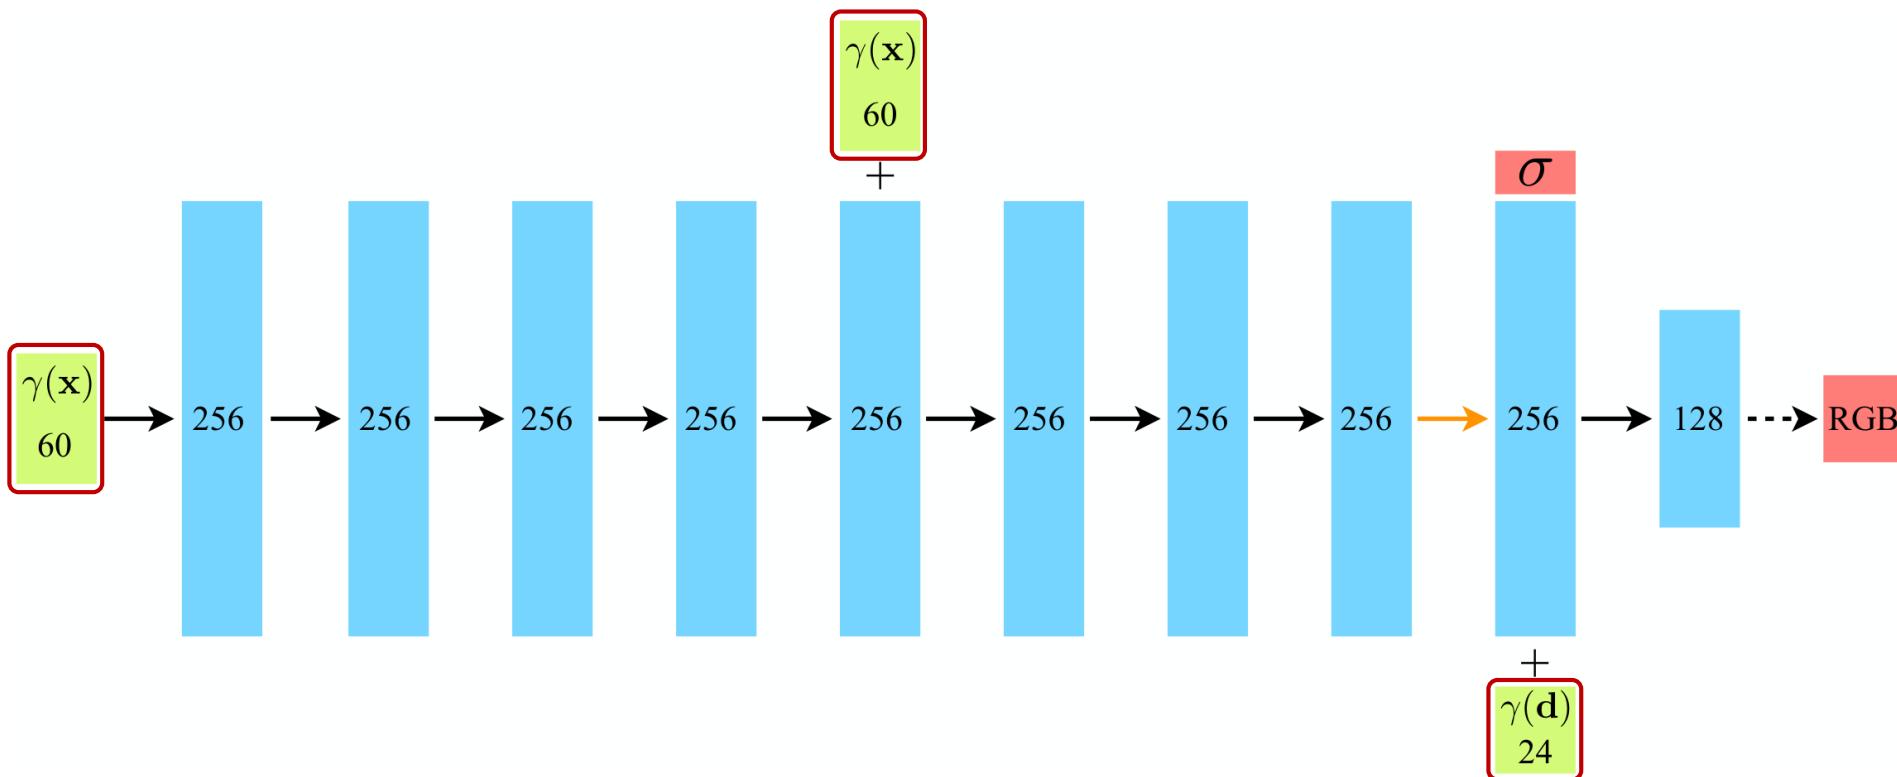
- **Inference:** Synthesizing a 2D image with a *new* viewpoint
  - Input: A new camera viewpoint ( $R_n, t_n$ )
- **Training:** Learning the 3D scene
  - Input: Images with their 3D viewpoints ( $R_j, t_j$ )
    - Note) 3D viewpoints can be retrieved by SfM (e.g. COLMAP).
  - Loss function: **Rendering loss (MSE)** between the *input* and *synthesized* images at each ( $R_j, t_j$ )
    - The *synthesized* images are generated by the neural **volumetric rendering**.



# NeRF (Neural Radiance Field; 2020)

- Network: **11 fully-connected (shortly FC) layers**

- Input:** Spatial location  $\mathbf{x} = (x, y, z)$  and viewing direction  $\mathbf{d} = (d_x, d_y, d_z)$  on a unit sphere
    - Q) What is a function  $\gamma$ ? Why are the input dimensions 60 and 24?
  - Output:** RGB color (RGB) and density ( $\sigma$ )

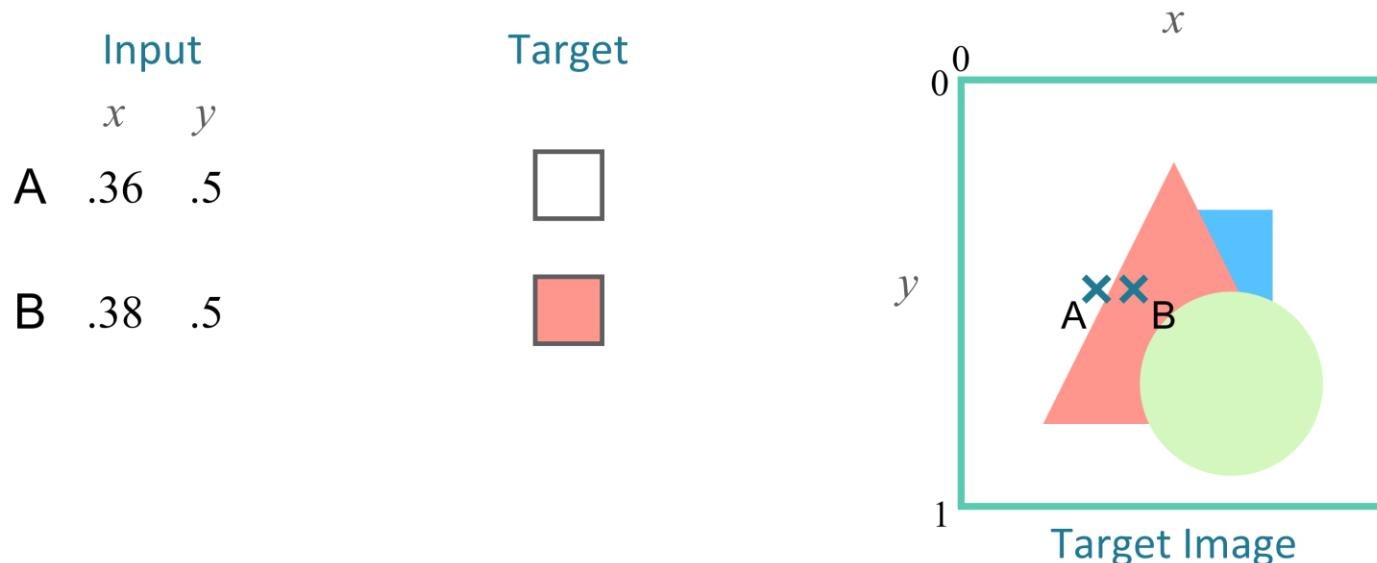


# NeRF (Neural Radiance Field; 2020)

- Key idea: **Positional encoding** transforms a real number  $p \rightarrow 2L$ -dimensional real numbers.

$$\gamma(p) = (\sin(2^0\pi p), \cos(2^0\pi p), \dots, \sin(2^{L-1}\pi p), \cos(2^{L-1}\pi p))$$

- Q) Why? 3.1415 and 3.1414 may generate similar values.

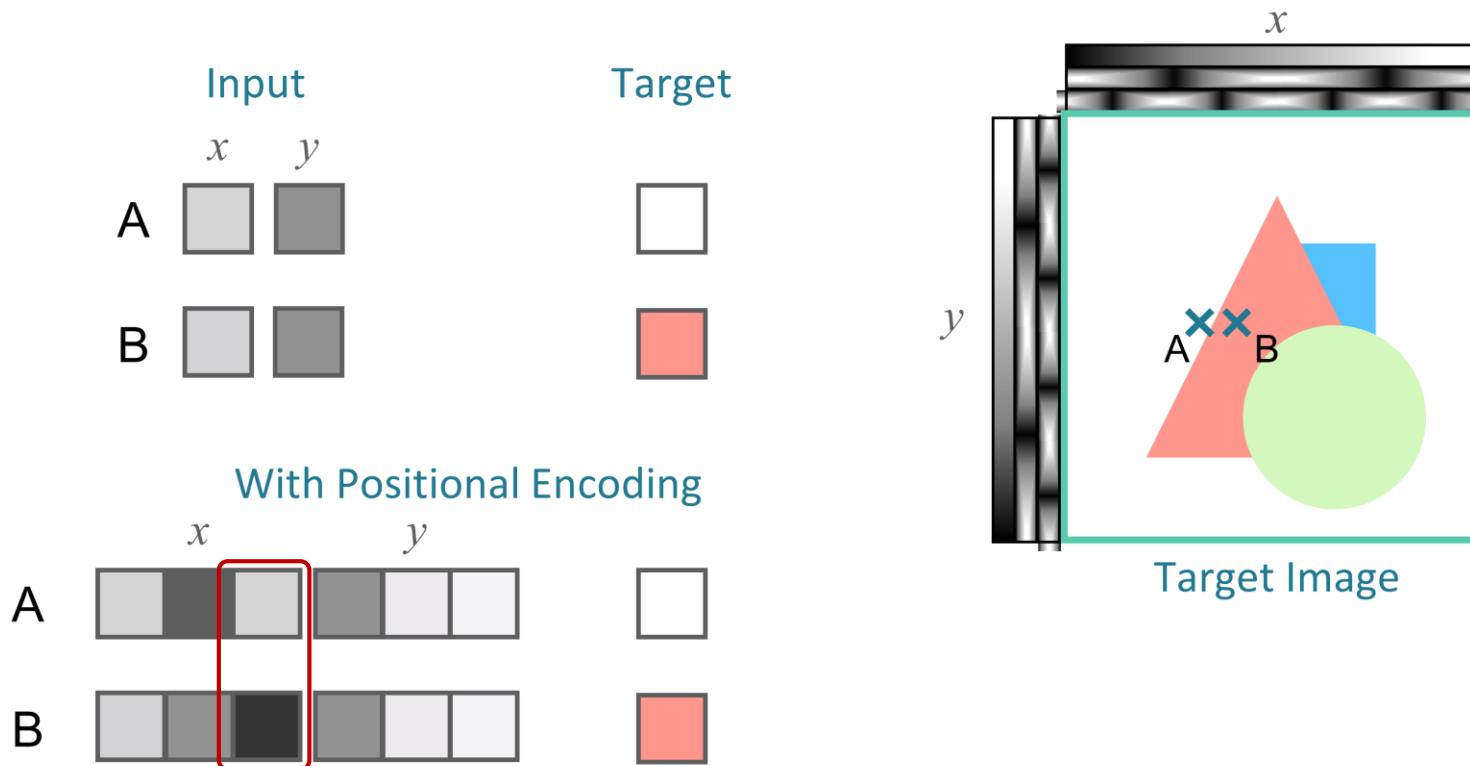


# NeRF (Neural Radiance Field; 2020)

- Key idea: **Positional encoding** transforms a real number  $p \rightarrow 2L$ -dimensional real numbers.

$$\gamma(p) = (\sin(2^0\pi p), \cos(2^0\pi p), \dots, \sin(2^{L-1}\pi p), \cos(2^{L-1}\pi p))$$

- Q) Why? 3.1415 and 3.1414 may generate similar values.

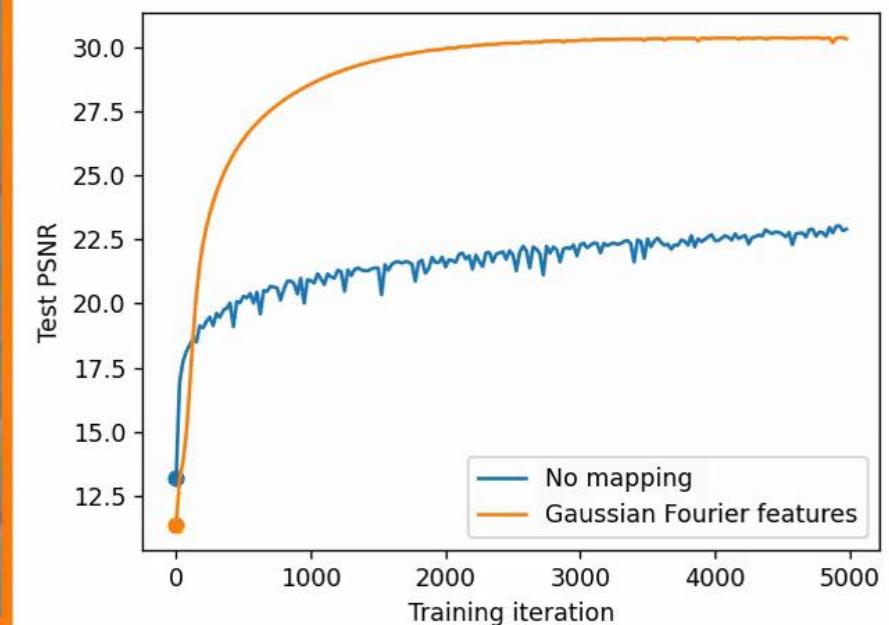
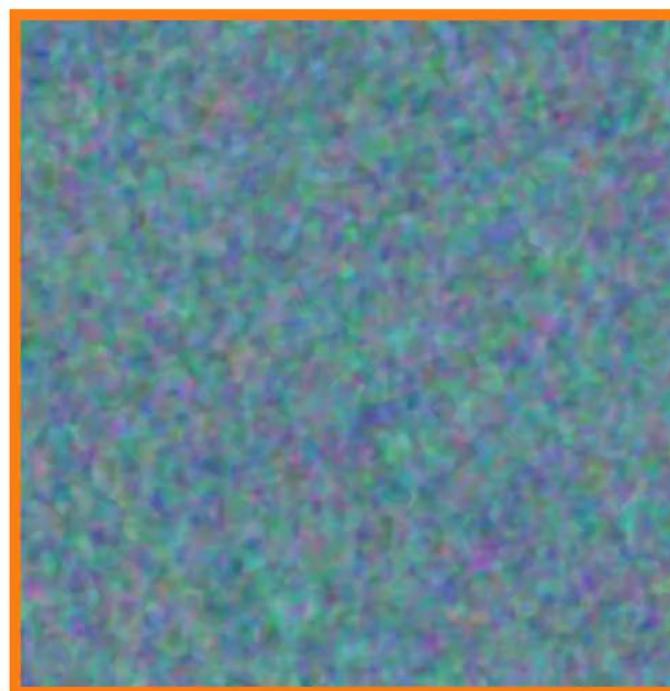


# NeRF (Neural Radiance Field; 2020)

- Key idea: **Positional encoding** transforms a real number  $p \rightarrow 2L$ -dimensional real numbers.

$$\gamma(p) = (\sin(2^0\pi p), \cos(2^0\pi p), \dots, \sin(2^{L-1}\pi p), \cos(2^{L-1}\pi p))$$

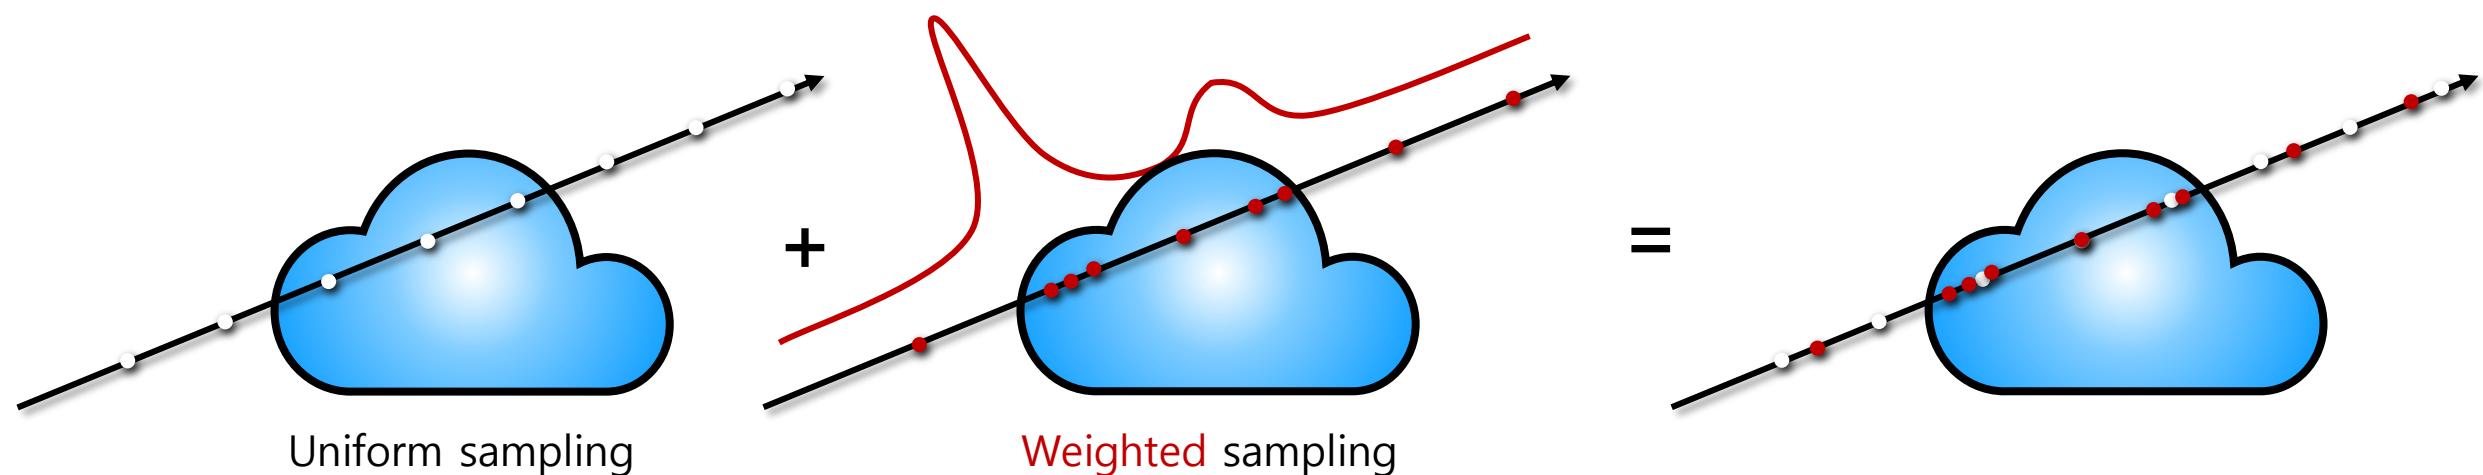
- Q) Why? 3.1415 and 3.1414 may generate similar values.
- A) Positional encoding **highlights** not only large values but also **small fraction numbers**.
- Note)  $L = 10$  for  $\gamma(\mathbf{x})$  and  $L = 4$  for  $\gamma(\mathbf{d})$



# NeRF (Neural Radiance Field; 2020)

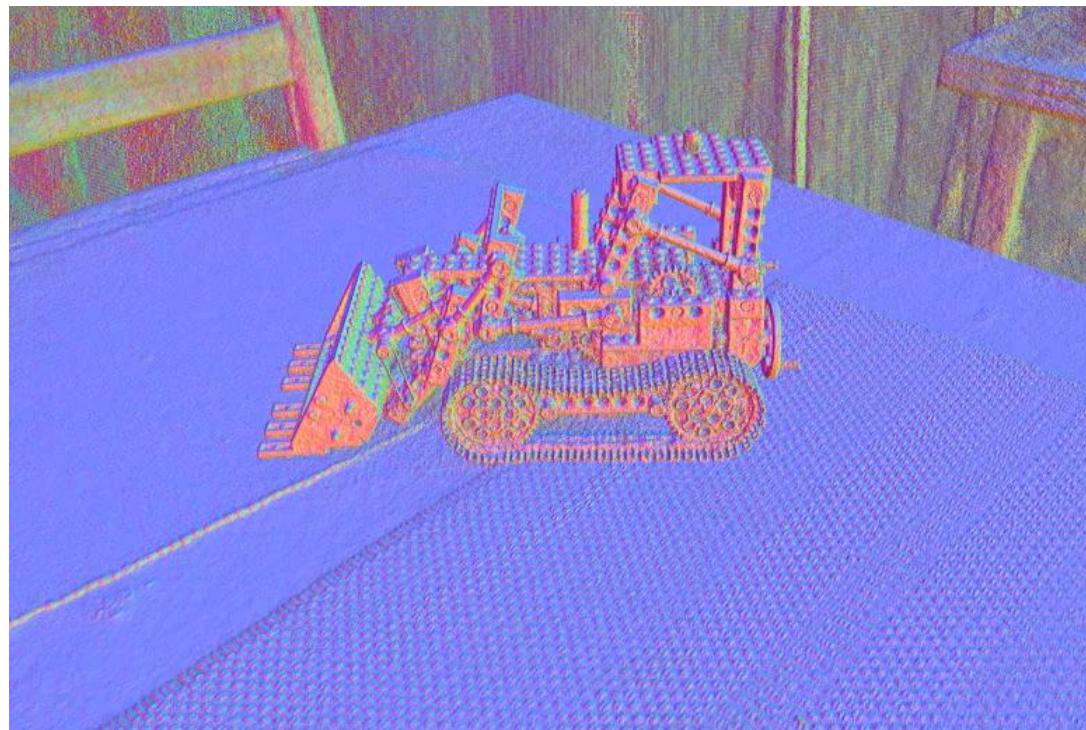
- Key idea: **Hierarchical volume sampling** selects points according to uniform and non-uniform weights.
  - Q) How to assign **non-uniform weights**?

$$C(\mathbf{r}) = \sum_{i=1}^N T_i \alpha_i \mathbf{c}_i = \sum_{i=1}^N w_i \mathbf{c}_i \rightarrow \hat{w}_i = \frac{w_i}{\sum w_i}$$



# NeRF (Neural Radiance Field; 2020)

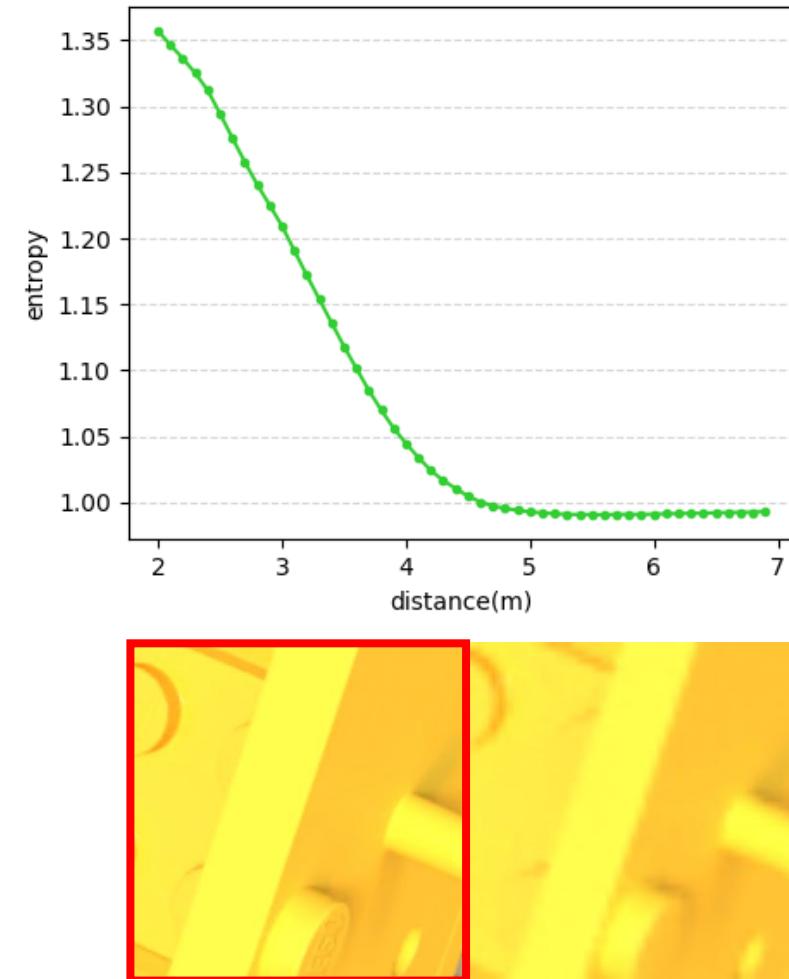
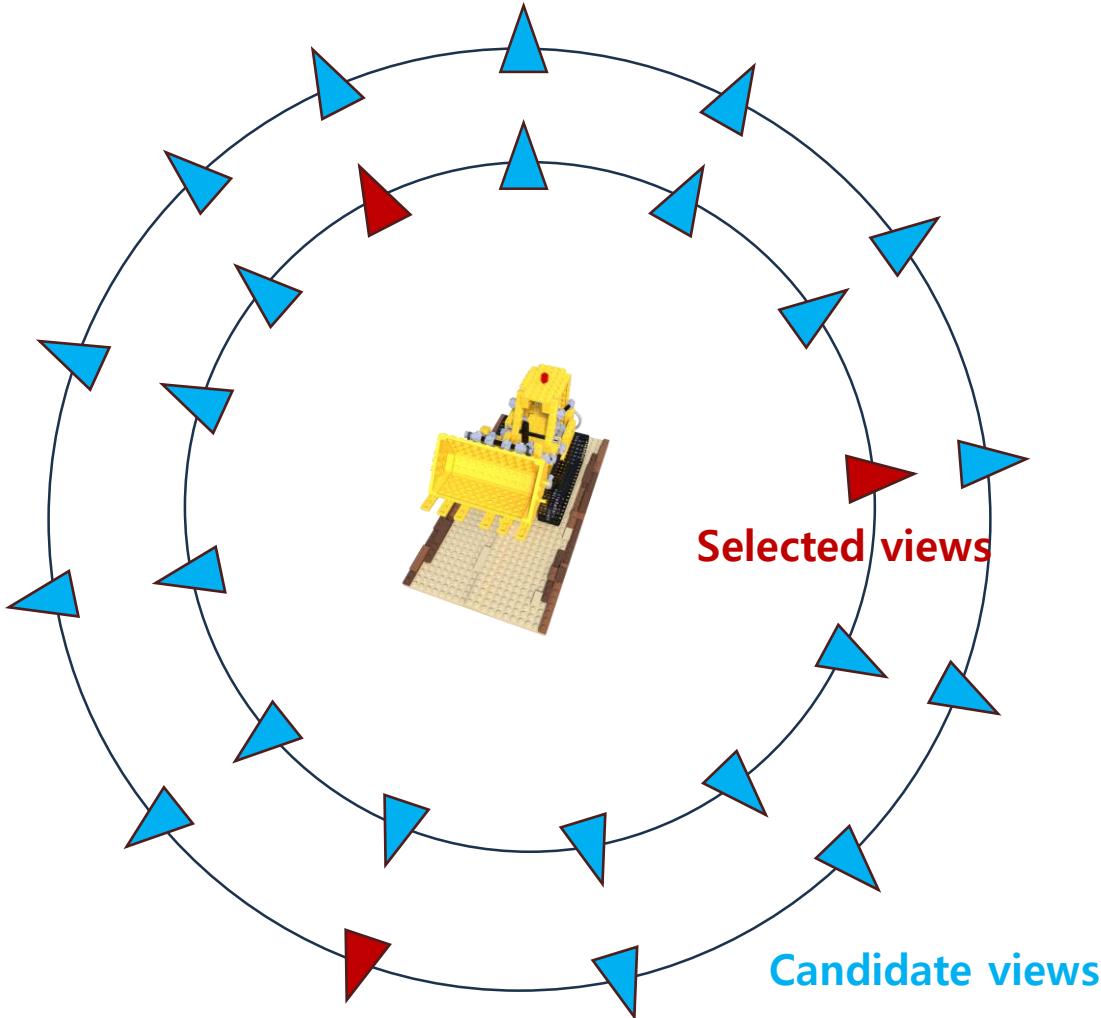
- **Inference:** Synthesizing a **2D image** with a *new* viewpoint
  - Input: A new camera viewpoint ( $R_n, t_n$ )
  - Neural volumetric rendering
- **Inference:** Retrieval of a **3D model** from density values
  - e.g. Normal vectors from analytic gradient of density





# ETRI-3DV Project: Next-Best-View Selection for Complete 3D Reconstruction

- 3D representation: **NeRF**
- Uncertainty measure: **Entropy** (interpreting density  $\rho$  as probability) with **distance-based regularization**





# ETRI-3DV Project: Next-Best-View Selection for Complete 3D Reconstruction

- Evaluation: [Lee et al., RA-L, 2023] vs. Proposed Method

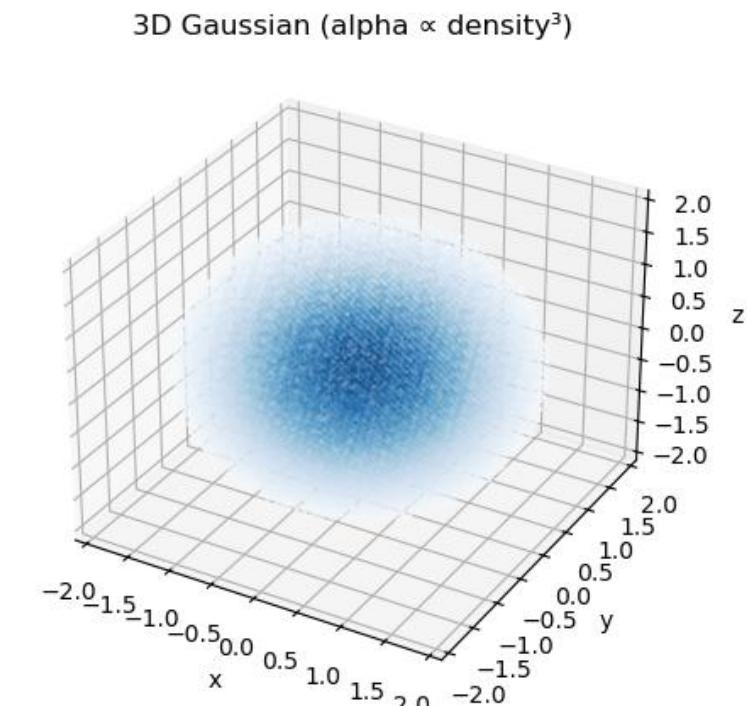
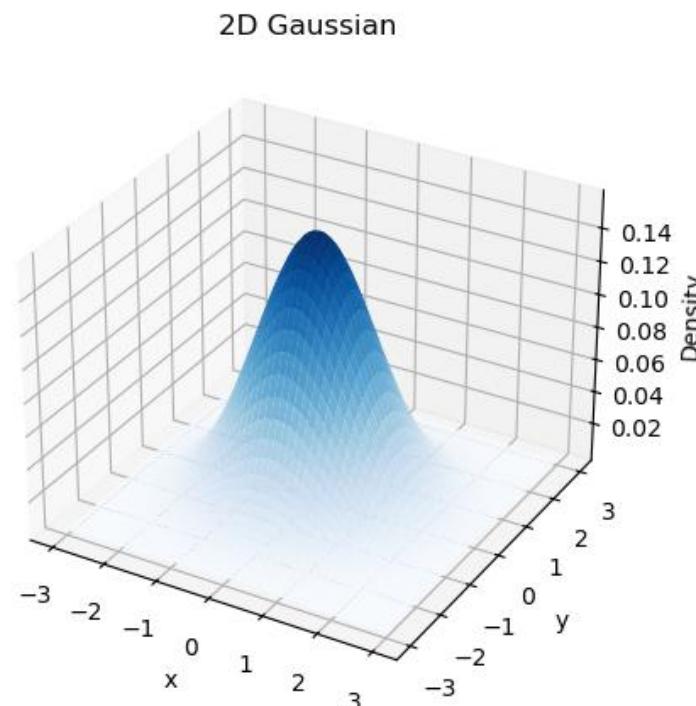
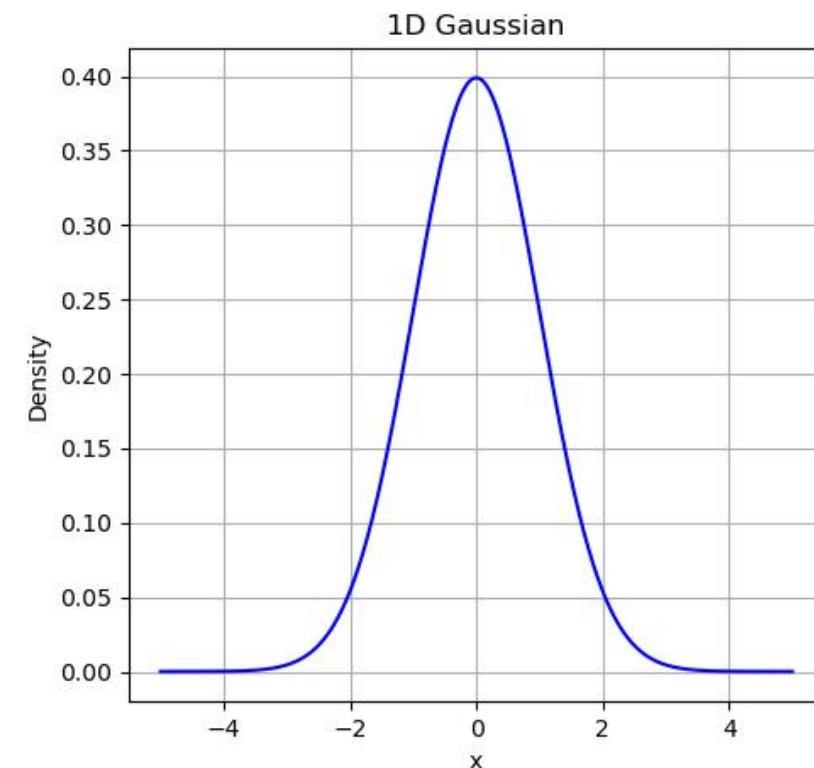


Objects	The Number of NBV Images					
	20		40		60	
	Entropy [2]	Ours	Entropy [2]	Ours	Entropy [2]	Ours
Chair	0.1538	<b>0.0561</b>	0.0892	<b>0.0239</b>	0.0536	<b>0.0173</b>
Hotdog	0.1026	<b>0.0950</b>	0.0746	<b>0.0426</b>	0.0327	<b>0.0258</b>
Mic	0.2521	<b>0.0108</b>	0.0134	<b>0.0080</b>	0.0065	<b>0.0041</b>

[Table. 1] Chamfer distance (Note: lower ( $\downarrow$ ) is better.) of the original entropy-based method [2] and our proposed entropy-based methods with distance regularization on three different object datasets

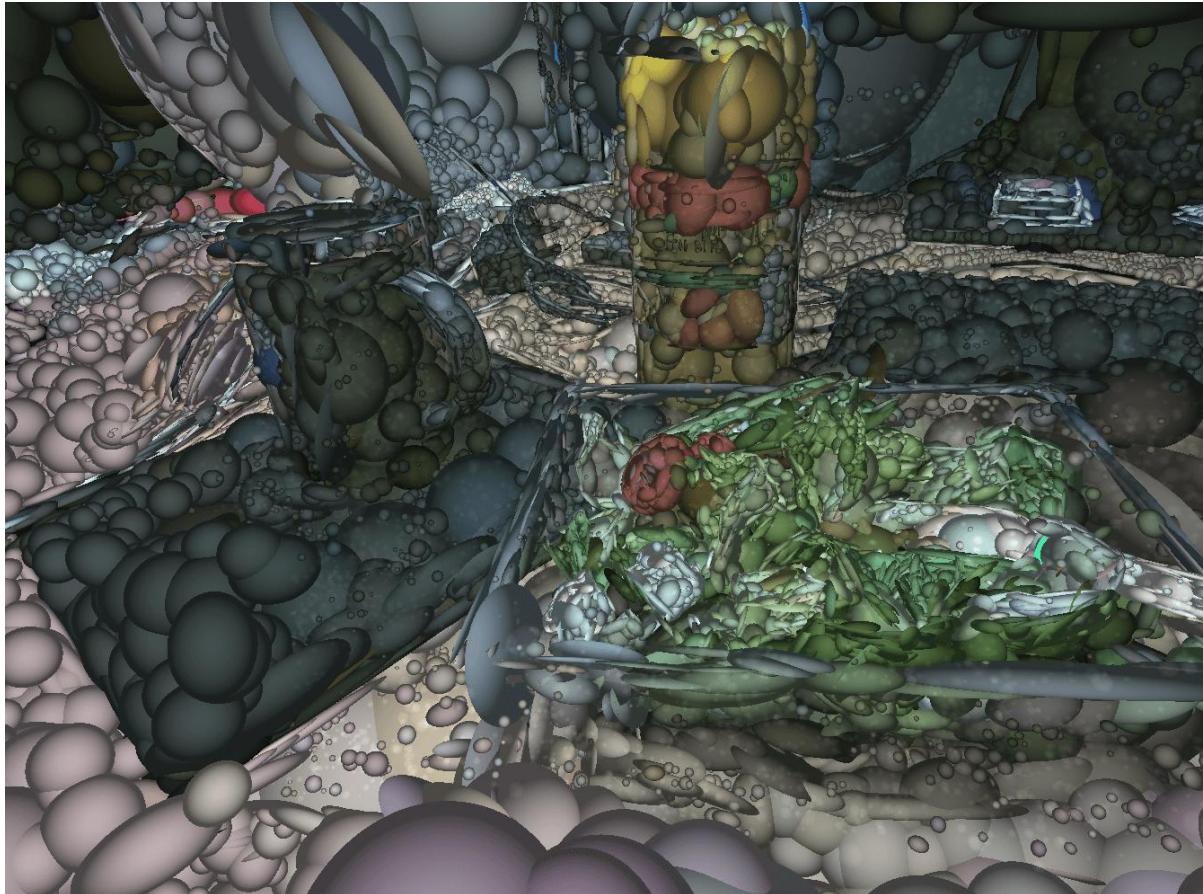
# 3D Gaussian Splatting (3DGS; 2023)

- 3DGS is a *explicit* 3D representation with a **collection of 3D Gaussians** for fast and high-quality rendering.
  - **Gaussians:**  $g(\mathbf{x}) = \exp\left(-\frac{1}{2} \mathbf{x}^\top \Sigma^{-1} \mathbf{x}\right)$



# 3D Gaussian Splatting (3DGS; 2023)

- 3DGS is a *explicit* 3D representation with a **collection of 3D Gaussians** for fast and high-quality rendering.
  - 3D Gaussians ~ 3D ellipsoids (an extension of *point cloud*)



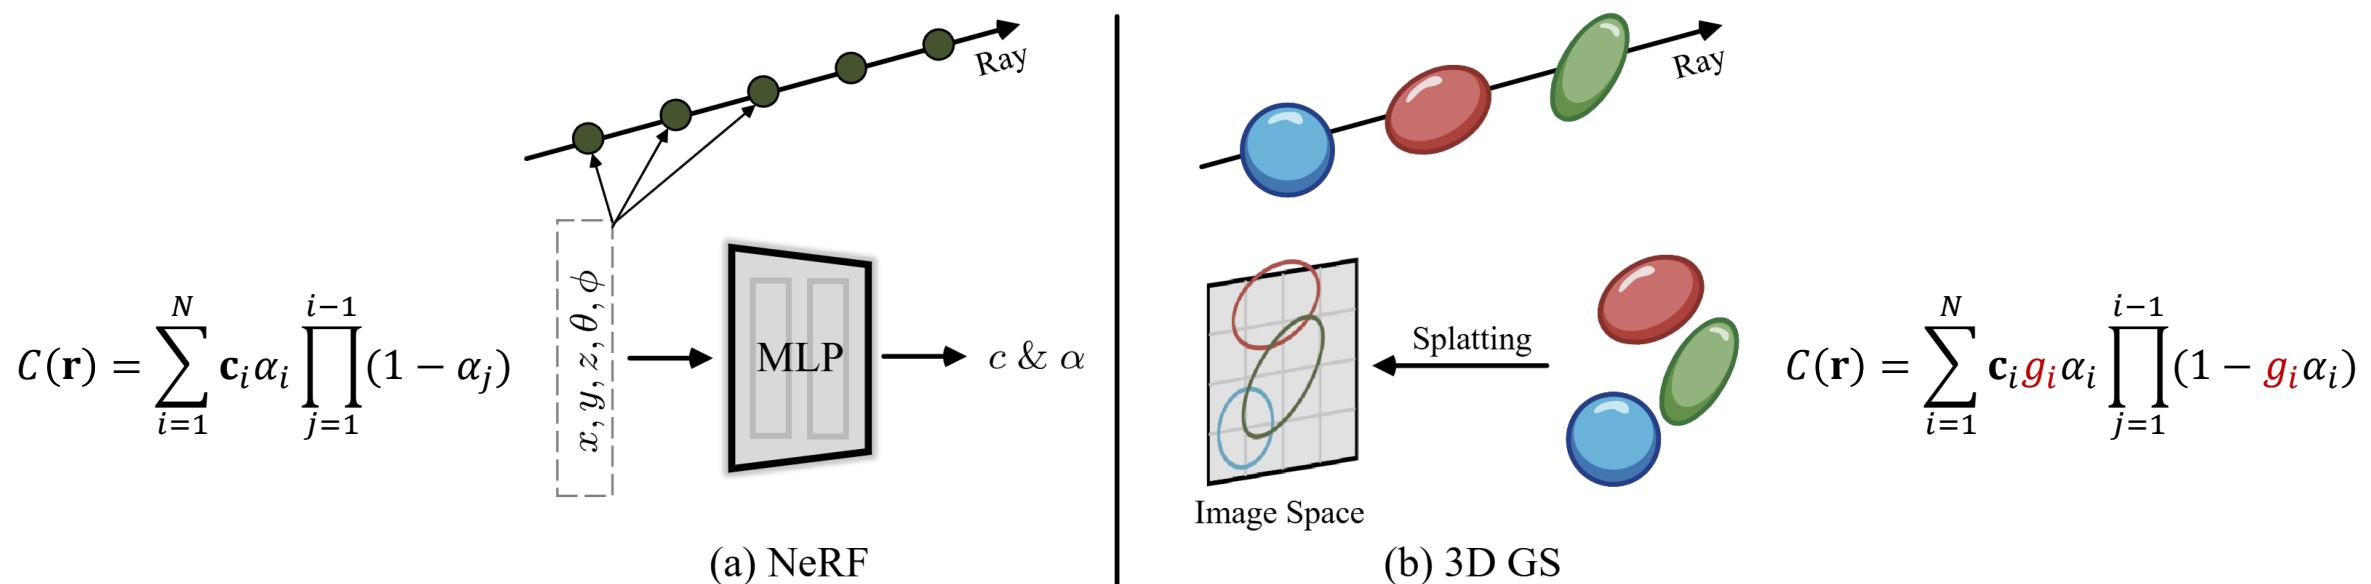
3D Gaussian visualization



Its rasterized image

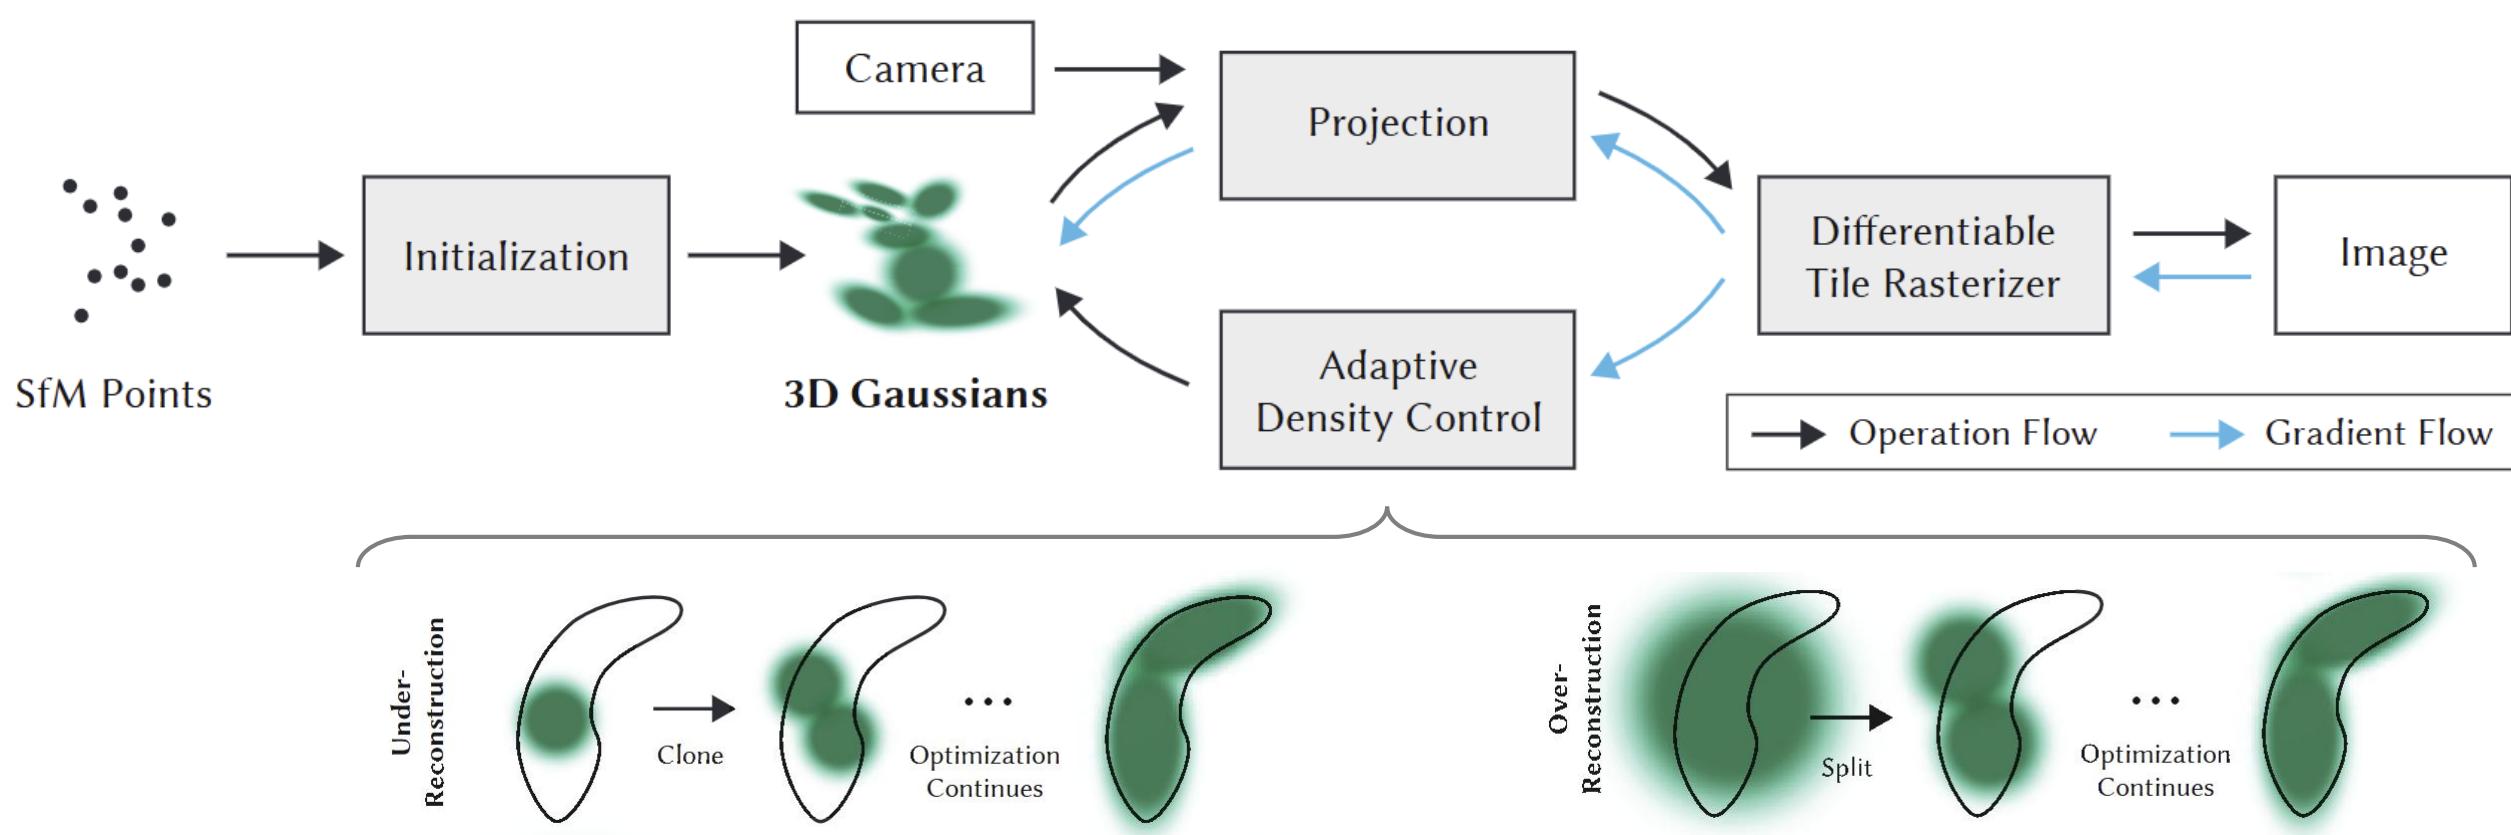
# 3D Gaussian Splatting (3DGS; 2023)

- 3DGS is a *explicit* 3D representation with a **collection of 3D Gaussians** for fast and high-quality rendering.
  - **3D Gaussians:**  $g(\mathbf{x}) = \exp(-\frac{1}{2} \mathbf{x}^\top \Sigma^{-1} \mathbf{x})$ 
    - Mean: 3D position  $\mathbf{x} = [x, y, z]^\top$
    - Covariance: 3D distribution  $\Sigma = RSS^\top R^\top$  ( $S$ : scale (anisotropic),  $R$ : rotation matrix)
    - Extra: Color (RGB), opacity ( $\alpha$ ), (optionally) view-dependent appearance (via spherical harmonics)
  - **Rendering** (Rasterization) = **Sort** (based on depth) + **Projection** (a.k.a. splatting) + (alpha-weighted) **Blending**



# 3D Gaussian Splatting (3DGS; 2023)

- 3DGS is a *explicit* 3D representation with a **collection of 3D Gaussians** for fast and high-quality rendering.
  - **Rendering** (→ Operation Flow) = **Sort** (based on depth) + **Projection** (a.k.a. splatting) + (alpha-weighted) **Blending**
  - **Training** (→ Operation Flow + ← Gradient Flow)
    - Loss function:  $\mathcal{L} = (1 - \lambda)\mathcal{L}_1 + \lambda\mathcal{L}_{\text{D-SSIM}}$



# 3D Gaussian Splatting (3DGS; 2023)

- 3DGS is a *explicit* 3D representation with a **collection of 3D Gaussians** for fast and high-quality rendering.



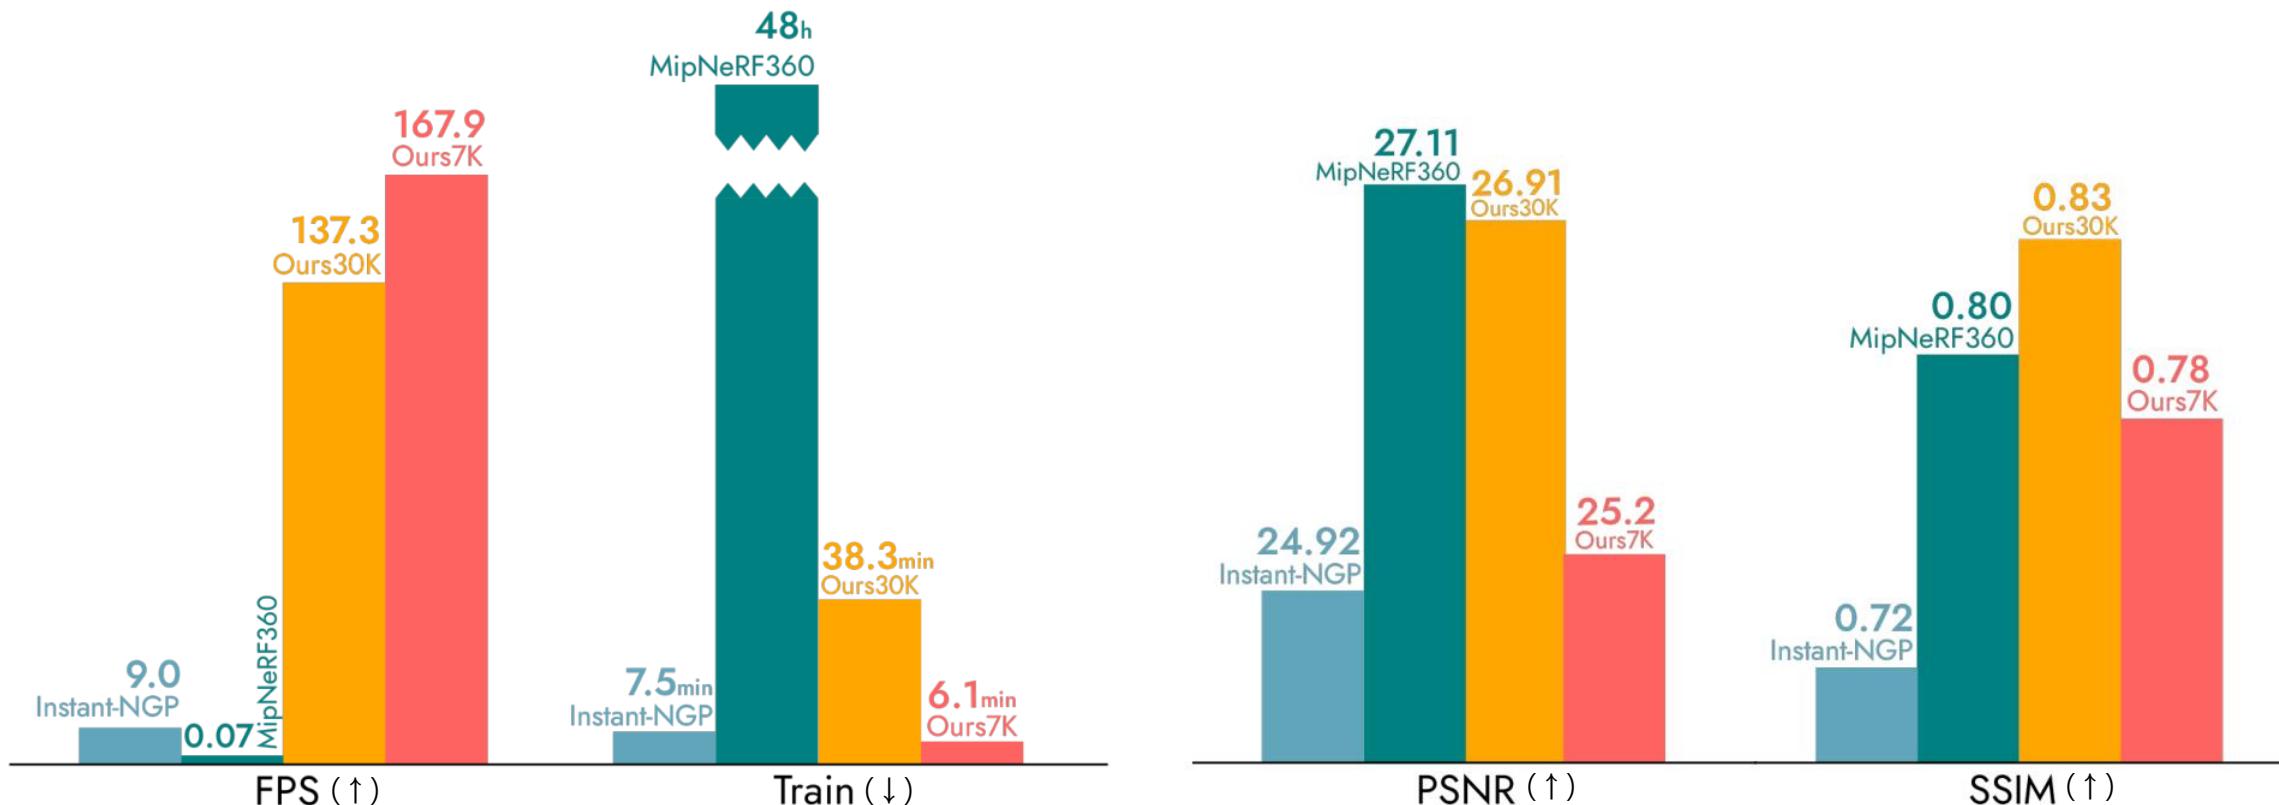
Mip-NeRF360



3DGS

# 3D Gaussian Splatting (3DGS; 2023)

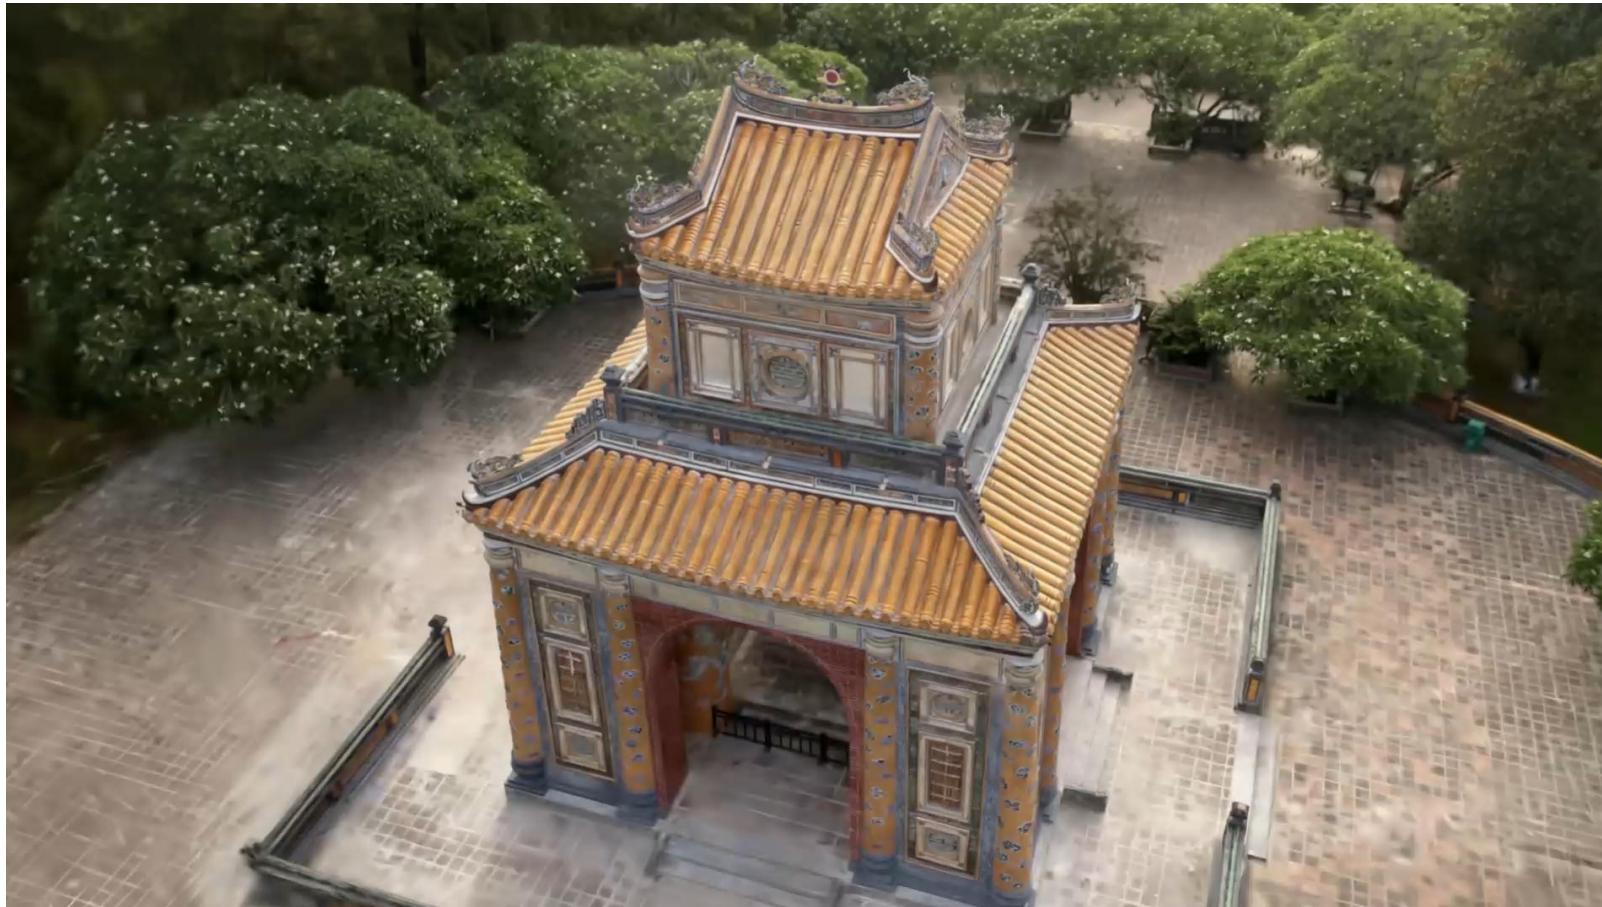
- **3DGS** is a *explicit* 3D representation with a **collection of 3D Gaussians** for fast and high-quality rendering.
  - NeRF suffers from slow training and rendering.
  - Evaluation @ Full MipNeRF360 Dataset + 2 Tanks and Temples + 2 Deep Blending



- Memory usages (↓): Instant-NGP (15-50MB), MipNeRF360 (8.6MB), 3DGS (350-700MB @ 3-6M of Gaussians)

# 3D Gaussian Splatting (2023)

- Applications: Real-time 3D engines, 3D capture tools, VFX, ...
  - Unreal Plugin: [XVERSE 3D-GS UE Plugin](#)
  - Unity Plugin: [Gaussian Splatting Playground in Unity](#), [SplatVFX](#)
  - Polycam: [Gaussian Splat Tool \(community works\)](#)



# 3D Gaussian Splatting (2023)

- Applications: SfM, Visual SLAM/odometry, ...
  - SfM, Visual SLAM/odometry: [COLMAP-Free 3DGS](#), [Gaussian Splatting SLAM](#), [GS-SLAM](#)

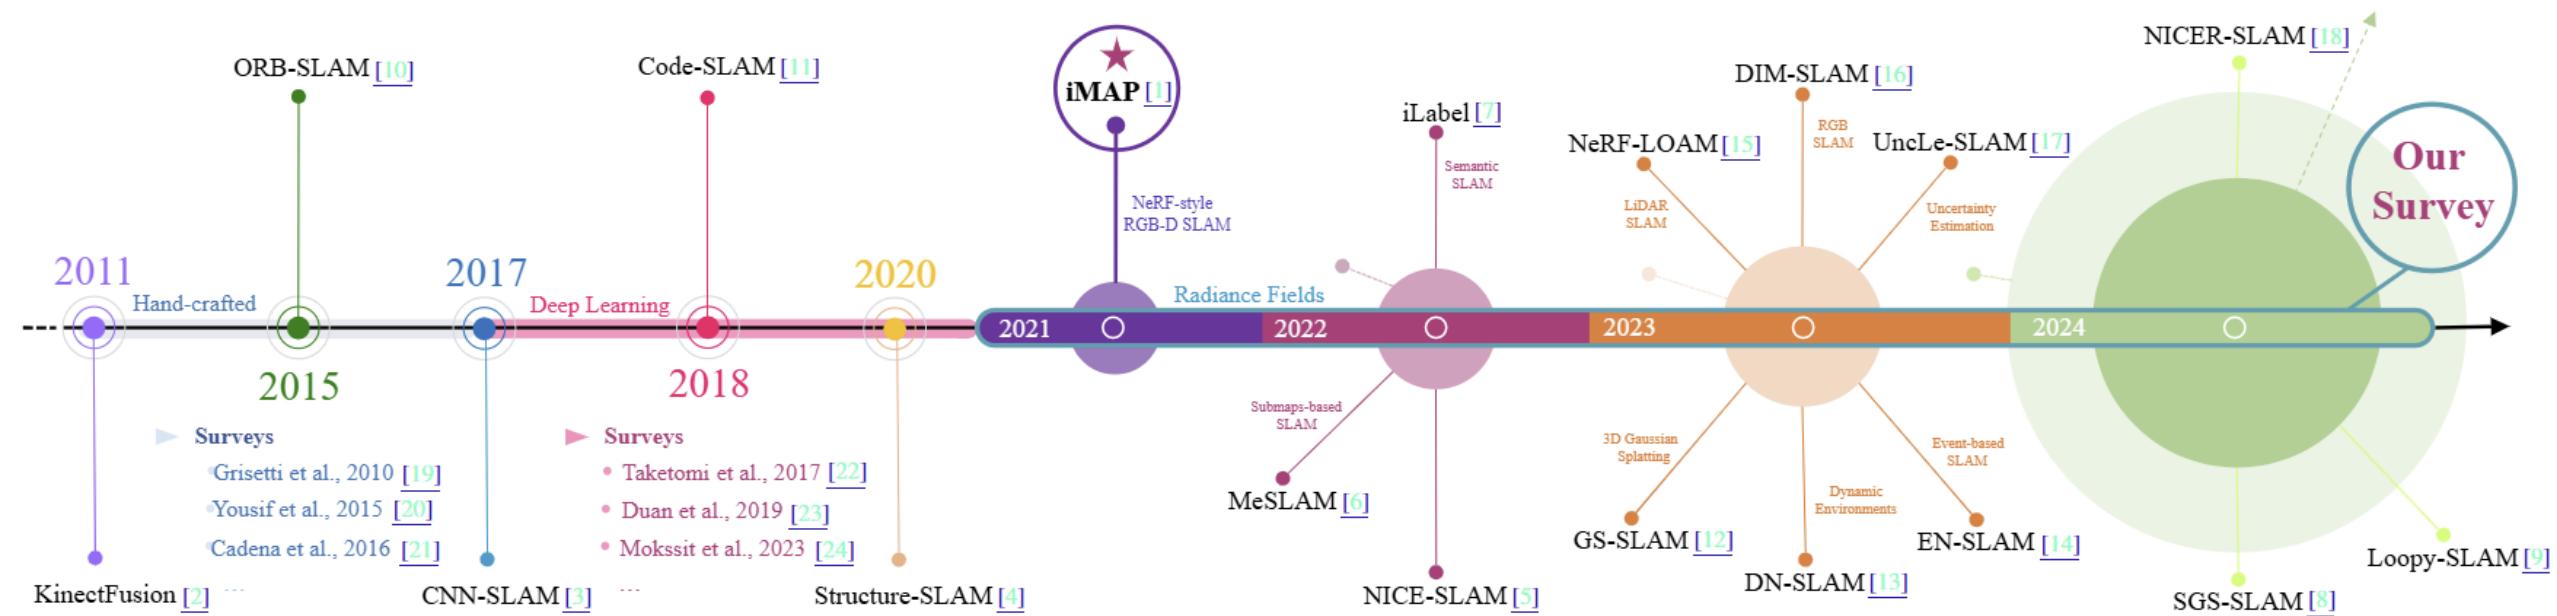


Fig. 1: **Timeline SLAM Evolution.** This timeline shows the evolution from hand-crafted to deep learning SLAM, with key surveys marking both periods. A significant shift occurs in 2021 with iMap [1], introducing radiance-field-based approaches. Circle sizes on the right indicate yearly publication volumes, with 2024's outer circle projecting increased interest in NeRF and 3DGS-based SLAM.

# 3D Gaussian Splatting (2023): Getting Started with gsplat

- Example) [gsplat](#) with the [relief dataset](#)

- Reconstruction results

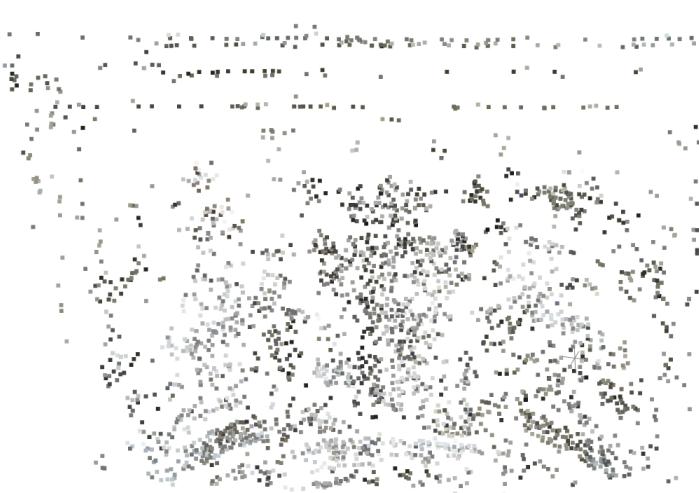
**Sparse Reconstruction (SfM)**



**Dense Reconstruction (MVS)**

or

**3D Gaussian Reconstruction**



# of points: 2,889



# of points: 336,223



Viewer: [Viser](#)

# of Gaussians: 221,767

# Summary

- **3D Representations:** How to represent 3D scenes and models

- Classical representations: Voxel, point cloud, polygon mesh, signed distance field (SDF)
- Neural Radiance Field (**NeRF**): 11 fully-connected (shortly FC) layers
  - Key idea: Neural volumetric rendering, positional encoding, hierarchical volume sampling  
→ High-quality view synthesis for *continuous* 3D scenes, but *too slow* rendering and training time
- 3D Gaussian Splatting (**3DGGS**): A collection of 3D Gaussians
  - Key idea: Volumetric splatting, (tricky) adaptive density control  
→ Faster and more high-quality, but *more memory* consumption
  - Applications: Real-time 3D engines, 3D capture tools, VFX, SfM, visual SLAM/odometry, ...